



# LECTURE 1: INTRODUCTION

University of Washington, Seattle

Fall 2025



# OUTLINE

## **Part 1: Welcome to EEP 596**

- Instruction team
- Goal of the course
- Instruction format
- Course materials
- Instruction components
- Schedule and Canvas page
- Syllabus and Grading

## **Part 2: Deep Learning and Neural Networks**

- Definition(s) of Deep Learning
- Model and data
- Fixed vs Learned model
- Classical machine learning methods
- Artificial neural network and Brain

## **Part 3: Regression**

- Regression problem
- Polynomial fit: Linear regression
- Improving linear regression



PART 1:

WELCOME TO EEP 596

# INSTRUCTION TEAM



**Instructor:**

Jimin Kim  
UW NeuroAI Lab  
jk55@uw.edu

## **Research Interests**

Computational Neuroscience  
and AI




**TA:**

Yang Zheng  
UW NeuroAI Lab  
jk55@uw.edu

## **Research Interests**

Neural network efficiency and  
Generative AI



**TA:**

Urvashi Taki  
utaki@uw.edu

## **Research Interests**

Deep learning for visual data  
and LLM

# WHAT IS THIS COURSE ABOUT?

**Fundamental concepts, skills and applications** of deep learning.

Learn **fundamental principles** behind deep learning **architecture** and **training**.

**Survey models** leading up to current state of the art methods.

Apply models to **real-world problems and datasets**: Labeled, Visual, Time-series etc.

Comprehensive introductory course with **heavy emphasis on practical aspects**.

Uses **Python 3 and PyTorch** as main programming tools.

# INSTRUCTION (Tentative)

(630pm – 640pm) Introduction and announcements

(640pm – 700pm) Canvas quiz discussion

**(700pm – 820pm) Lecture: Theory**

(820pm – 830pm) Break

**(830pm – 910pm) Lecture: Practice** (with live coding examples)

(910pm – 920pm) Lab assignment introduction

**(920pm – 950pm) In-person Lab session**

# COURSE MATERIALS

## **Weekly Assignment page** (Canvas)

- 1) Lecture slides/recording (.pdf, Panopto video)
- 2) Lab slides/recording (.pdf, Panopto video)
- 3) In-lab examples (.ipynb)
- 4) Lab report starter templates (.ipynb)

## **Additional Resources** (Canvas)

Deep Learning (Ian Goodfellow)

Neural Networks and Deep Learning (Michael Nelson)

PyTorch Github Tutorials (<https://github.com/yunjey/pytorch-tutorial>)

# INSTRUCTION COMPONENTS

- |                     |   |                               |
|---------------------|---|-------------------------------|
| 1. Lecture (Theory) | → | Theoretical Concepts          |
| 2. Canvas quiz      | → | Theoretical Concepts Feedback |
| 3. Lecture (Lab)    | → | Practical Concepts            |
| 4. Code examples    | → | Code Implementations          |
| 5. Lab assignment   | → | Real-world Applications       |



# SCHEDULE

## **Weekly Instruction:**

Mon 6:30 PM – 9:50 PM (ECE 037)

## **Office Hour**

Jimin: TBA

Yang: TBA

Urvashi: TBA

# CANVAS PAGE

EE P 596 Au 24: Practical Introduction to Deep Learning  
Applications and Theory 




Image credit: Google DeepMind

## Welcome to EEP 596!

Welcome to EEP 596 “Practical Introduction to Deep Learning Applications and Theory”! This is a graduate level course aiming to provide **fundamental skills, concepts, and applications of deep learning** and neural networks for the investigation of complex datasets with heavy emphasis on hands-on practices.

**Home & Syllabus** tabs for  
Detailed Course Info

**Announcements** tab for  
important course  
announcements

**Discussions** tab for  
student – student  
student – instructor  
discussions

**Panopto** tab for  
Post-lecture recordings

# CANVAS ASSIGNMENT PAGE

## Lab 1 Report ↕

### LAB 1: PYTHON FOUNDATION

#### Associated lectures

(Theoretical Concepts)

Lecture 1 - Introduction (Panopto recording, Slides)

(Practice)

Lab 1 - Python Foundation (Panopto recording, Slides)

#### Lab materials

[Lab report guidelines](#) ↓ (IMPORTANT: READ THIS FIRST!)

[Lab 1 Examples notebook](#) ↓

ipynb file containing all the examples discussed in the lab videos. Use this to play with the examples yourself.

[Lab 1 Report Template](#) ↓

Zip file containing lab report template ipynb + exercise image files (You need image files to load problems in ipynb).

Unzip the file using windows or 7-zip.

Read **lab report guidelines**  
before starting your first  
assignment

Your lab report = Filled in  
report template notebook  
**(.ipynb)**

# SYLLABUS (Tentative)

## **Neural network fundamentals and feed-forward networks (09/29 – 10/13)**

(W1) Intro to ANNs/Regression | Setting up Python environment/Regression

(W2) Intro to Deep Learning and MLP | PyTorch Introduction/Classification with MLP

(W3) Convolutional Neural Networks | CNNs in PyTorch

## **Sequence models (10/20 – 10/27)**

(W4) Recurrent Neural Networks | RNNs in PyTorch

(W5) LSTM, GRU, Encoder-Decoder architectures | LSTM, GRU, Encoder-Decoder in PyTorch

## **Generative models (11/03 – 11/17)**

(W6) Attention and Transformers | Text processing with Transformers

(W7) Advanced CNNs | Image classification with ResNet

(W8) Generative Adversarial Networks | Image generation with GAN

## **Advanced models (11/24)**

(W9) Deep Reinforcement Learning | Control strategies with DeepRL

## **Final project (12/01 – 12/08)**

(W10) Project week | Selected Projects

(Finals week) Project presentation

# GRADING

(i): **Canvas Quiz (20%)** – individual, weekly

Evaluated on concept understanding and correctness of the responses

**Goal: Fundamental concepts feedback**

(ii): **Lab Assignment (40%)** – individual, weekly

Evaluated on code organization, documentation and task completion

**Goal: Implementation and training of Deep Learning models on real datasets**

(iii): **Final Project (40%)** – individual/team

Evaluated on project planning, presentation and code implementation

**Goal: Solve an original, real-world project of your choice using deep learning**

Q/A



# PART 2:

## WHAT IS DEEP LEARNING?



# Definition(s) of Deep Learning

**IBM:** Subset of machine learning that uses multilayered neural networks, called deep neural networks, to **simulate** the complex decision-making power of the human brain.

**AWS:** A method in artificial intelligence (AI) that teaches computers to **process** data in a way that is inspired by the human brain.

**Ian Goodfellow:** A form of machine learning that enables computers to **learn** from experience and understand the world in terms of a hierarchy of concepts.


**Yann LeCun:** Constructing networks of parameterized functional modules & **training** them from examples using gradient-based optimization.

**Geoffrey Hinton:** An approach to machine learning that involves **computational models** with multiple processing layers to learn representations of data with multiple levels of abstraction



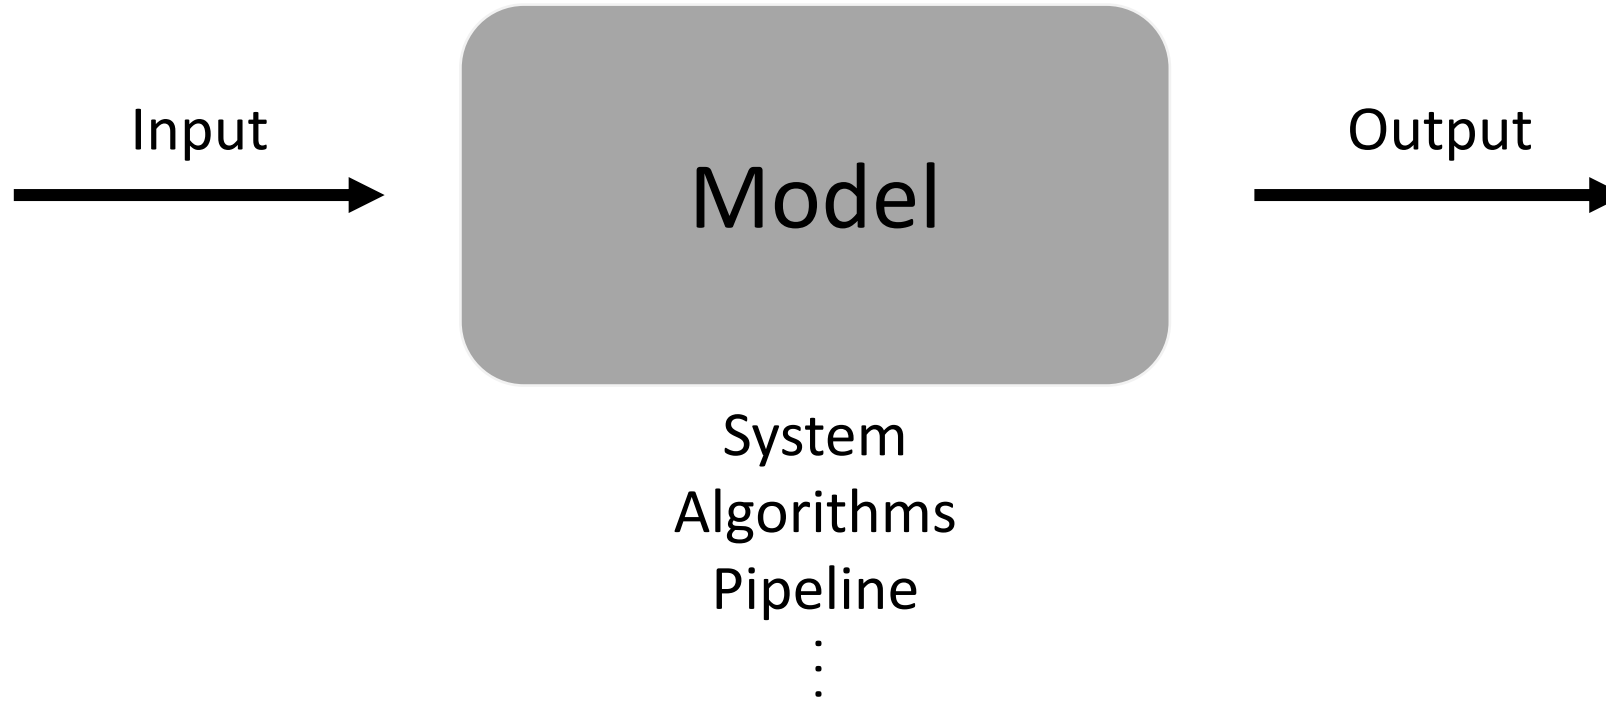


# Deep Learning Foundation: Model and Data



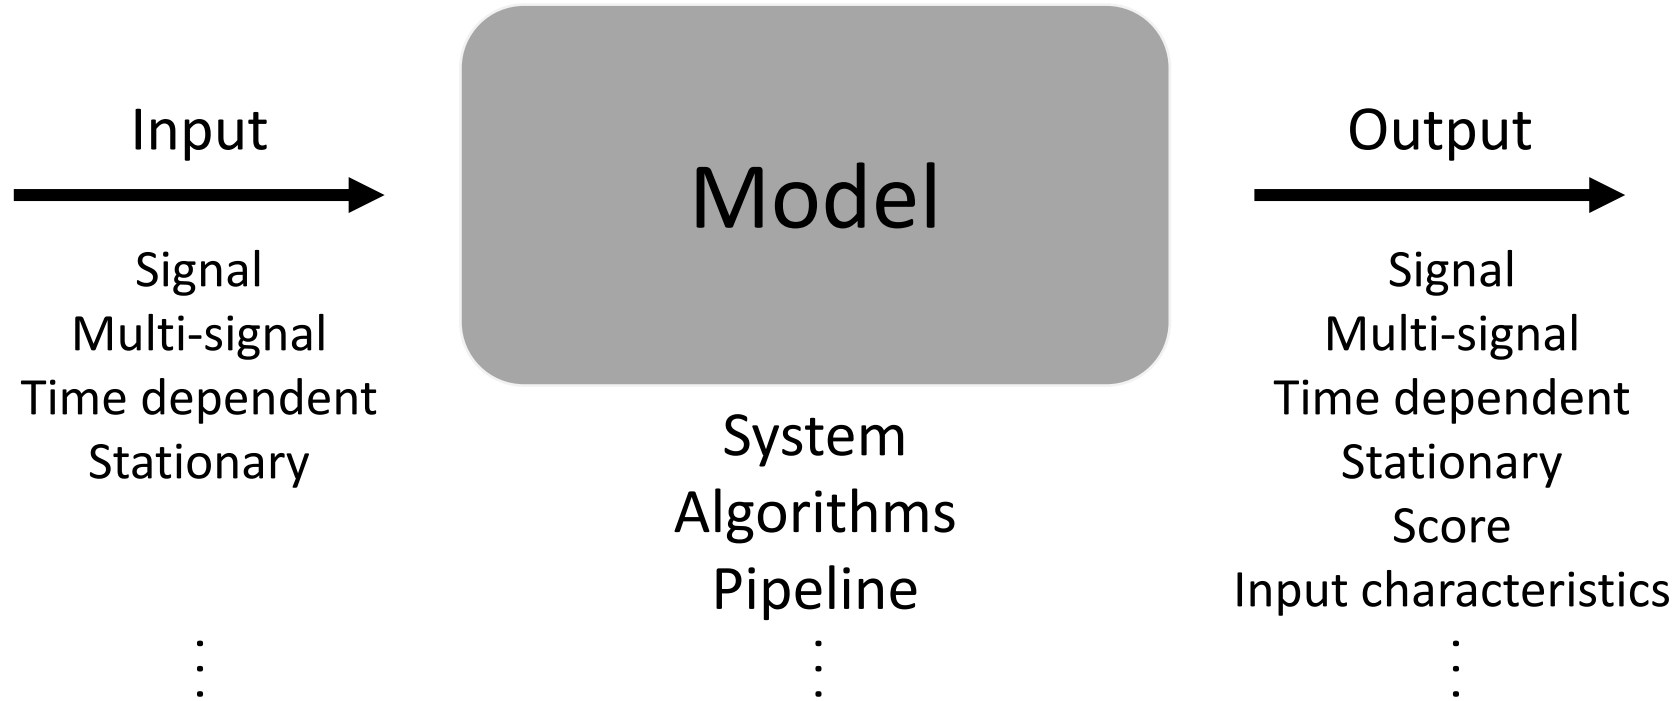


# Model and Data



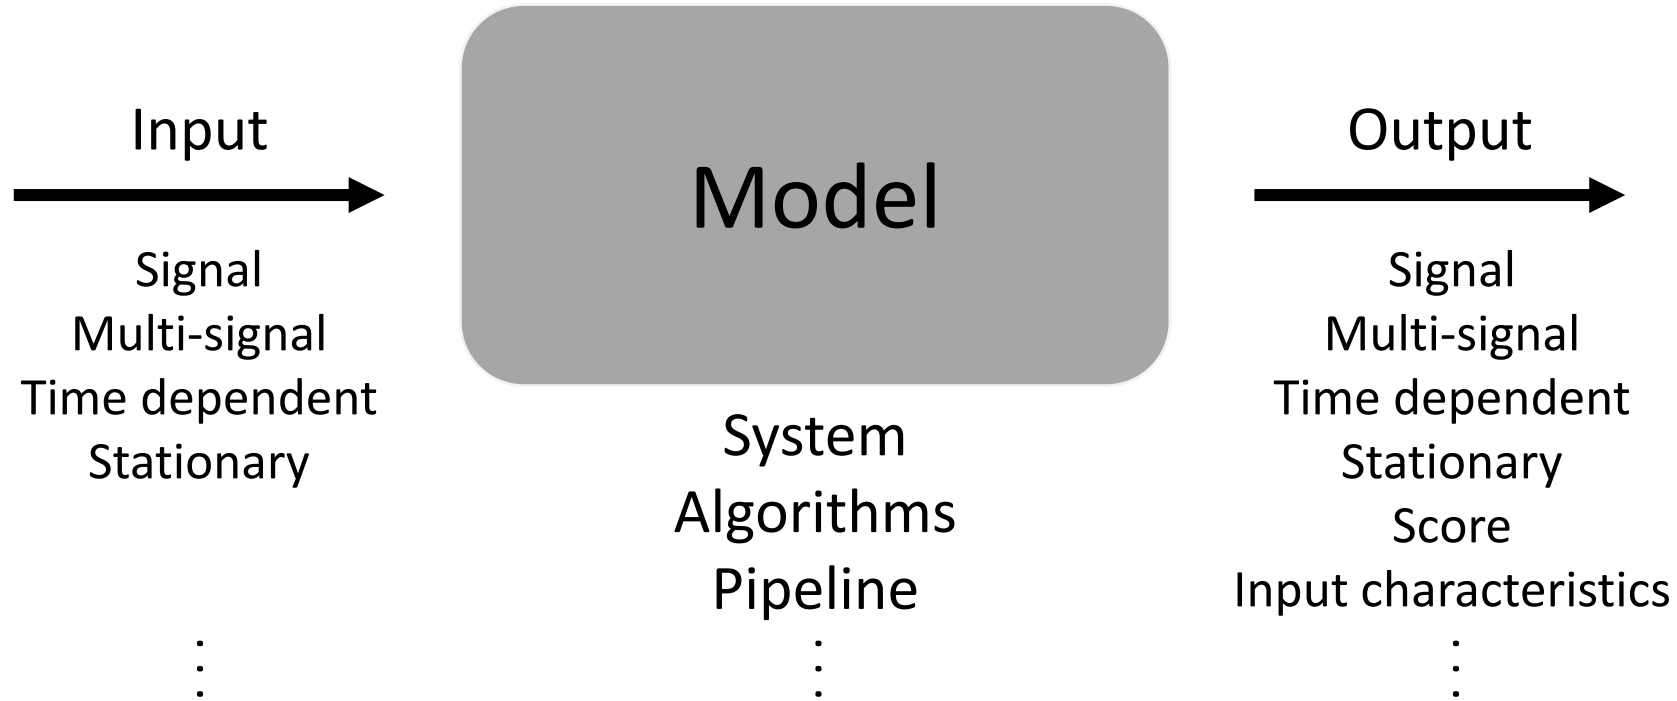


# Model and Data





# Model and Data




**Model = Function of the input**



# Model examples

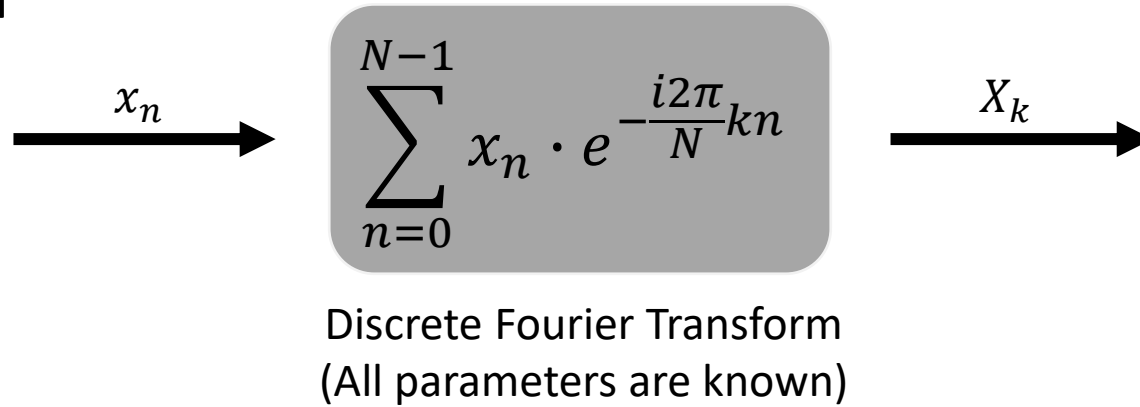
- LTI systems
  - Convolution, Fourier, Z-transform, Laplace transform
- Non-LTI systems
  - Non-linear transform
- Data fitting models
  - Linear regression
  - Polynomial regression
- Scoring models
- Classification models





# Fixed vs Learned model

Fixed model



Learned model


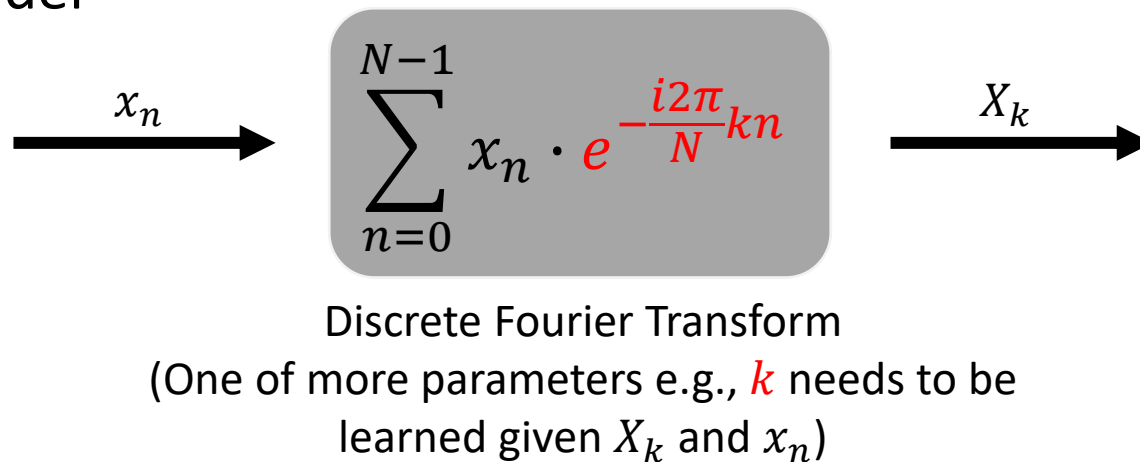
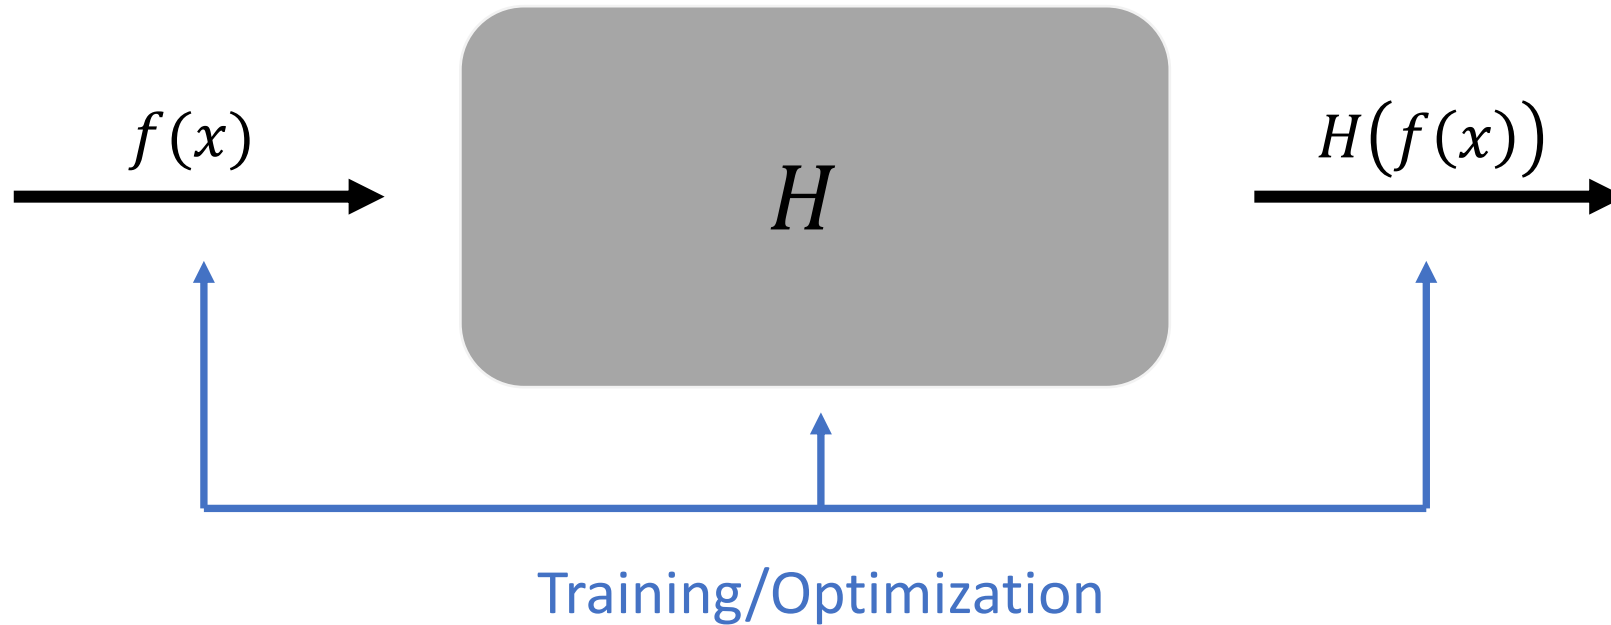


Image credit: All about circuits

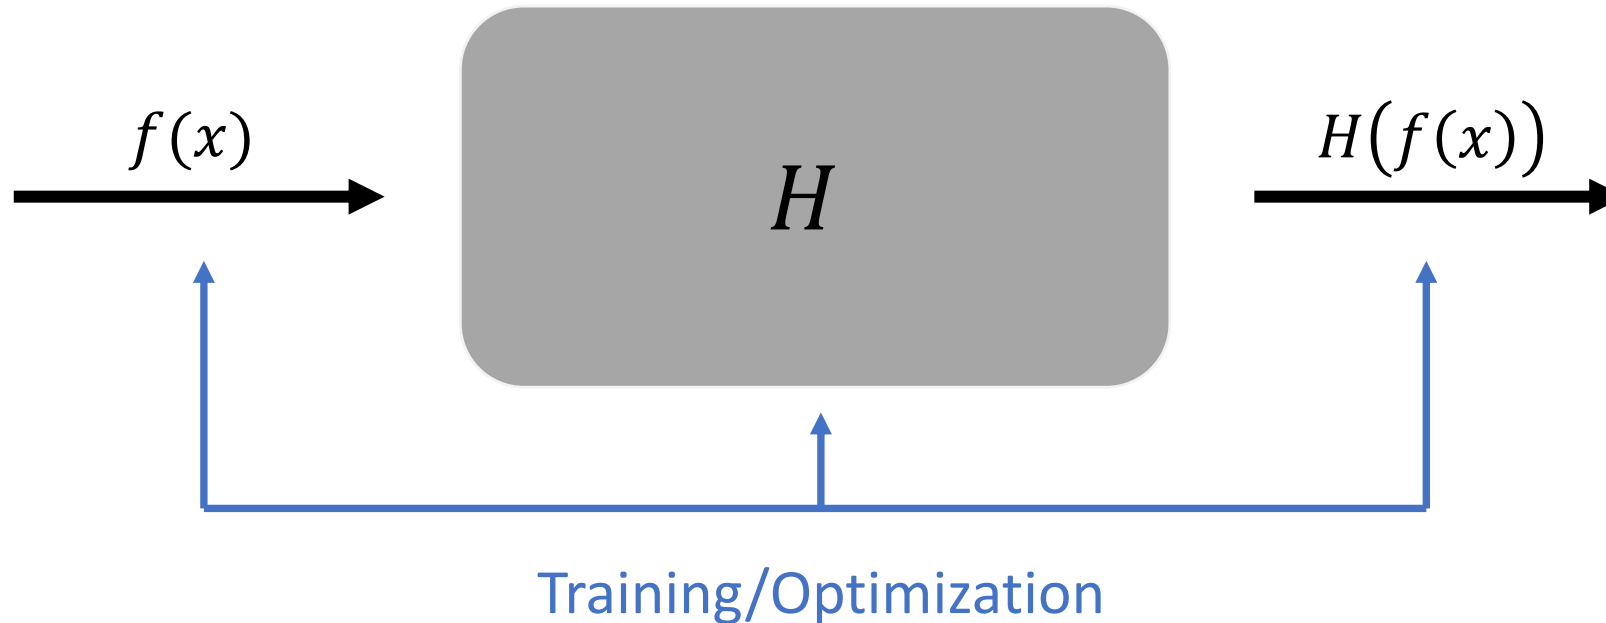


# Machine Learning = Learning an Optimal Model for Data





# Machine Learning = Learning an Optimal Model for Data




- $H$  expression unknown(Hard)
- $H$  expression fixed with known parameters (Easier)
- $H$  expression can be iteratively updated through Machine Learning algorithms (optimization, training)
- **Deep Learning  $\subset$  Machine Learning**






# Machine Learning = Learning an Optimal Model for Data

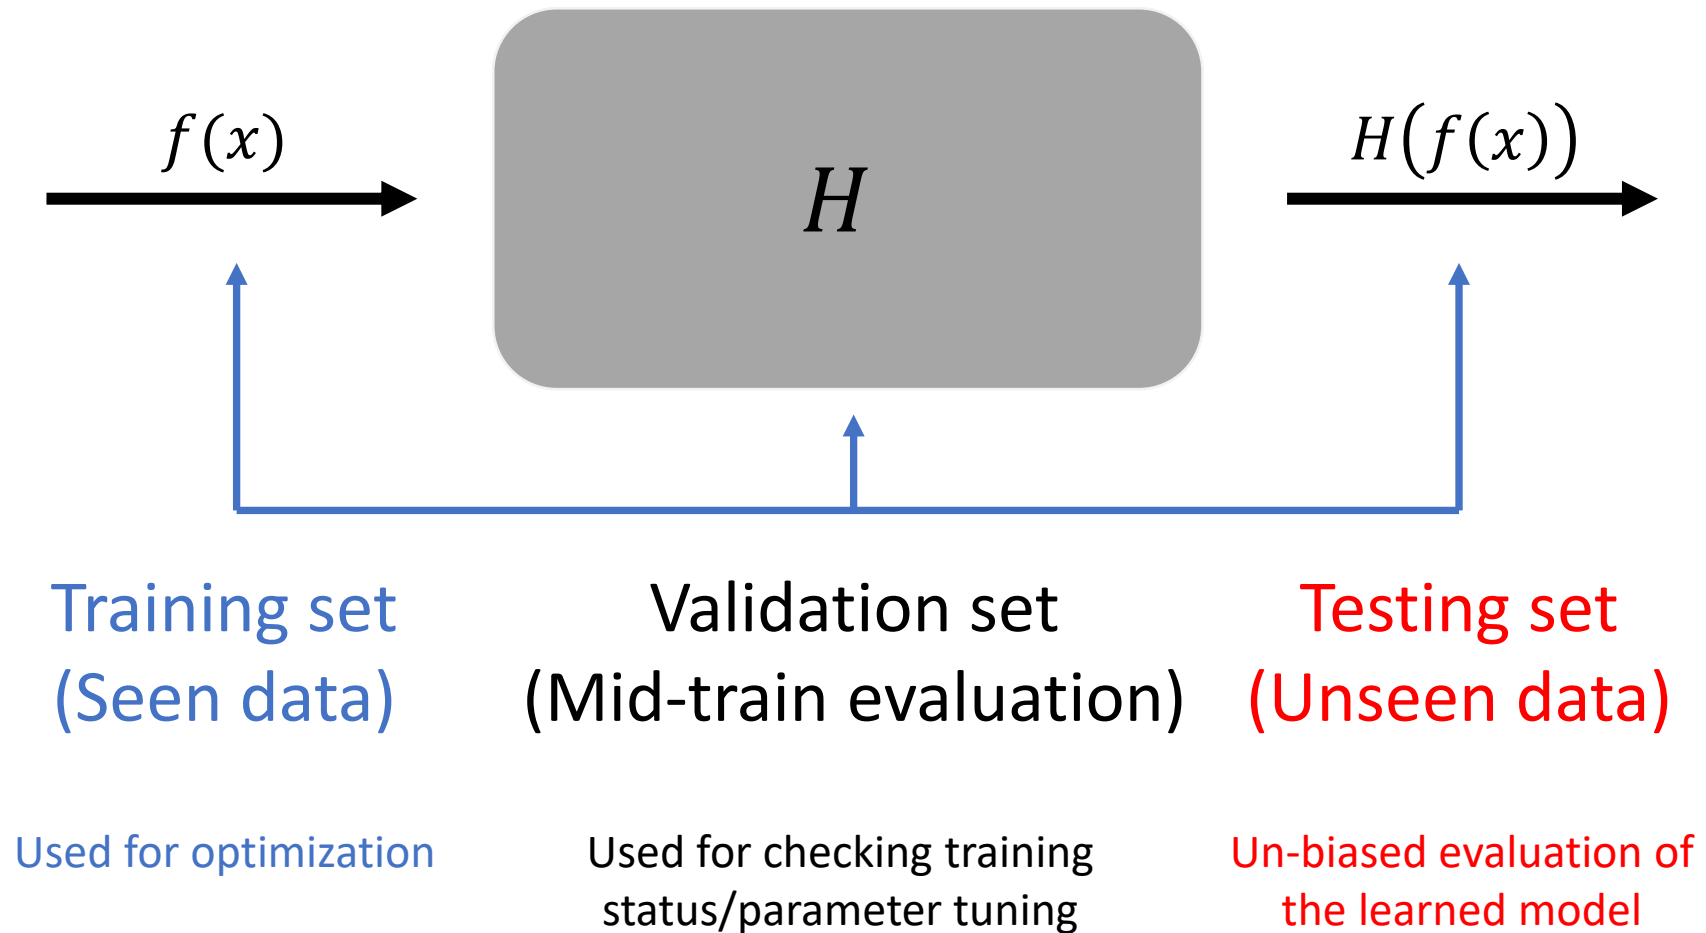


Model optimization is to convert a training set to a model which satisfies the training set – Ilya Sutskever

# Three Pillars of Machine Learning Training





# Three Pillars of Machine Learning Training







# Classical Machine Learning





# Classical Machine Learning




Feature extractions

- Round face
- Black eyes and nose
- Triangular ears



# Limitations of Classical Machine Learning




Model



1: Café Poodle  
0: Not Café Poodle



# Limitations of Classical Machine Learning




- Feature extractions are often done manually
- Not robust with data
- Hard to scale





# Neural Network Model




1: Café Poodle  
0: Not Café Poodle





# Neural Network Model




1: Café Poodle  
0: Not Café Poodle

- Handles (learns) feature extractions
- Robust with data
- Scalable





# Neural Network Model







# Lessons from brain







# Lessons from brain





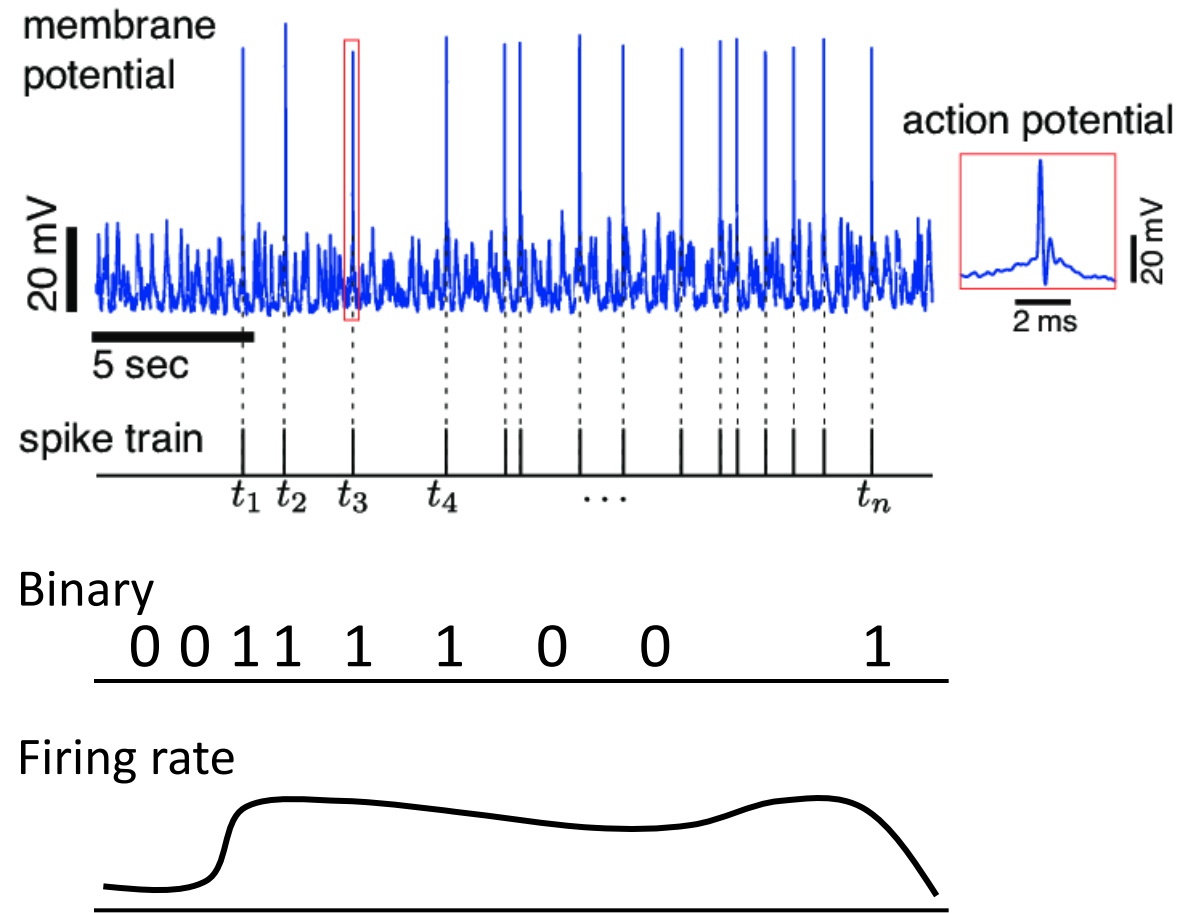

# Lessons from brain









# Lessons from brain








# Lessons from brain








# Lessons from brain





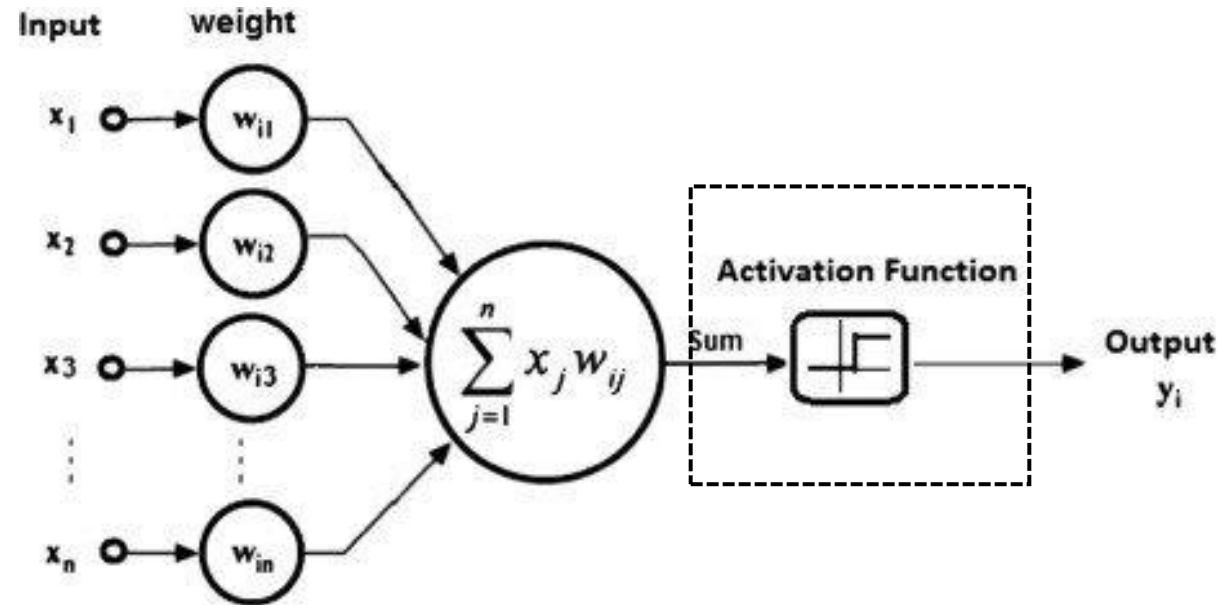


# Lessons from brain





# Mathematical Description of a Neuron



$$I = x_1 w_1 + \dots + x_n w_n$$

$$f(I)$$

$$x = f(I)$$


Integration

Activation

Output to other  
neurons



# Activation Functions in Neural Network



$$I = x_1 w_1 + \dots + x_n w_n$$

Integration

$$\sum_{i=1}^n x_i w_i + b$$

$$f(I)$$

Activation

$$f\left(\sum_{i=1}^n x_i w_i + b\right)$$

$$y = f(I)$$


Output to other neurons

$$y = f\left(\sum_{i=1}^n x_i w_i + b\right)$$

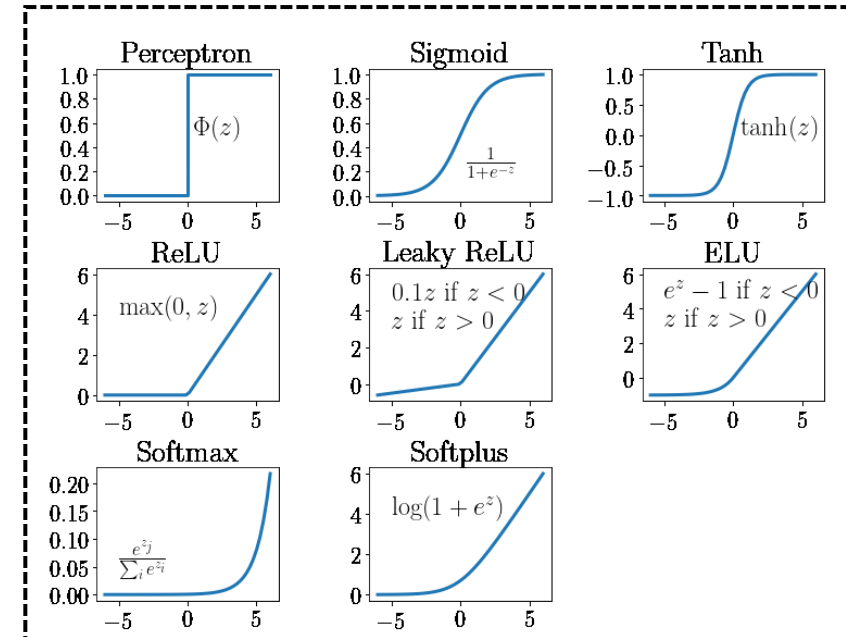


# Mathematical Description of a Neuron

Without non-linear activation,  
neural network becomes a  
linear model




Activation function introduces  
non-linearity to the model





# Neural Network Model



$$y = f\left(\sum_{i=1}^n x_i w_i + b\right)$$

Learn the **connection weights** and **biases** that optimizes the training set

# Biological Neural Networks are large and complex


Nematode *C. elegans*



302 neurons

*~7000 synapses*


Fruit fly



~150k neurons

*~70mil synapses*


Mouse



~70mil neurons

*~7 x 10<sup>8</sup> synapses*

Human




~10<sup>11</sup> neurons


*~10<sup>14</sup> synapses*

# Biological Neural Networks are large and complex

Cerebral cortex



Visual system




Biological neural networks are both Hierarchical and Recurrent (parallelism)





# Deep Neural Network as Brain Analogue










# Deep Learning Applications

Computer vision



Data processing



Text/speech  
generation

Image/video generation






# PART 3:

# Regression



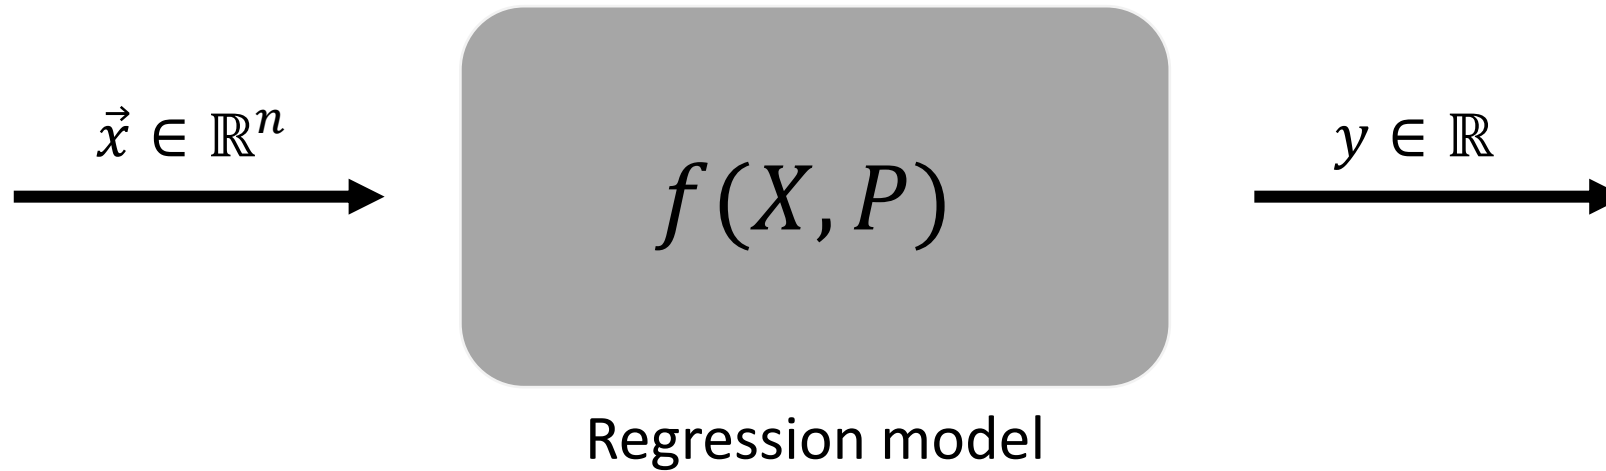
# Regression problem

$$Y = f(X, W)$$





# Regression problem




Linear regression (polynomial fit)

$$y(x) = p_0 + p_1x + \cdots + p_mx^m$$

$$y = \vec{p} \cdot \vec{x}$$



# Polynomial fit – Linear regression



$x = \text{input data}$

$y = \text{output data}$

$$\hat{y}(x) = p_0 + p_1x + \cdots + p_mx^m$$

Goal: Minimize  $e = y - \hat{y}$



# Polynomial fit – Linear regression

$$\hat{y}(x) = p_0 + p_1x + \cdots + p_mx^m$$

Goal: Minimize  $e = y - \hat{y}$



# Polynomial fit – Linear regression

$$\hat{y}(x) = p_0 + p_1x + \cdots + p_mx^m$$

Goal: Minimize  $e = y - \hat{y}$

$$(x_0, y_0), (x_1, y_1), \dots, (x_n, y_n)$$

n-given points

$$X = \begin{pmatrix} 1 & x_0 & x_0^2 & \cdots & x_0^m \\ 1 & x_1 & x_1^2 & \cdots & x_1^m \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_n & x_n^2 & \cdots & x_n^m \end{pmatrix} \begin{matrix} \vec{\tilde{x}}_0 \\ \\ \\ \vec{\tilde{x}}_n \end{matrix}$$

$$P = \begin{pmatrix} p_0 \\ p_1 \\ \vdots \\ p_m \end{pmatrix}$$



# Polynomial fit – Linear regression

$$\hat{y}(x) = p_0 + p_1x + \cdots + p_mx^m$$

Goal: Minimize  $e = y - \hat{y}$

$(x_0, y_0), (x_1, y_1), \dots, (x_n, y_n)$   
n-given points

$$X = \begin{pmatrix} 1 & x_0 & x_0^2 & \cdots & x_0^m \\ 1 & x_1 & x_1^2 & \cdots & x_1^m \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_n & x_n^2 & \cdots & x_n^m \end{pmatrix} \begin{matrix} \vec{\tilde{x}}_0 \\ \\ \\ \vec{\tilde{x}}_n \end{matrix} \quad P = \begin{pmatrix} p_0 \\ p_1 \\ \vdots \\ p_n \end{pmatrix}$$

Fit (“Learn”) the model:

Least squares error

**Cost**  $J = E_2 = \sum_{i=1}^n (\overbrace{\vec{p} \cdot \vec{\tilde{x}}_i}^{\hat{y}_i} - y_i)^2$   
**Total error**





# Polynomial fit – Linear regression

Find the parameters minimizing the error

$$J = E_2 = \sum_{i=1}^n (\vec{p} \cdot \vec{\tilde{x}}_i - y_i)^2$$

$$\vec{p}^* = \arg \min_{\vec{p}} J(\vec{p})$$



# Polynomial fit – Linear regression

Find the parameters minimizing the error

$$J = E_2 = \sum_{i=1}^n (\vec{p} \cdot \vec{\tilde{x}}_i - y_i)^2$$

$$\vec{p}^* = \arg \min_{\vec{p}} J(\vec{p})$$

$$\forall i : \quad \frac{\partial J}{\partial p_j} = 0 \qquad \frac{\partial J}{\partial p_j} = \sum_{i=1}^n 2(\vec{p} \cdot \vec{\tilde{x}}_i - y_i) \vec{\tilde{x}}_i$$

$$\vec{p}^* = (X^T X)^{-1} X^T \vec{y}$$



# Polynomial fit – Linear regression

Find the parameters minimizing the error

$$J = E_2 = \sum_{i=1}^n (\vec{p} \cdot \vec{\tilde{x}}_i - y_i)^2$$


$$\vec{p}^* = \arg \min_{\vec{p}} J(\vec{p})$$

$$\forall i : \quad \frac{\partial J}{\partial p_j} = 0 \qquad \frac{\partial J}{\partial p_j} = \sum_{i=1}^n 2(\vec{p} \cdot \vec{\tilde{x}}_i - y_i) \vec{\tilde{x}}_i$$

$$\vec{p}^* = (X^T X)^{-1} X^T \vec{y} \quad \text{Linear least square formula}$$




# Linear regression is not perfect



$$\hat{y}(x) = p_0 + p_1x + \cdots + p_mx^m$$



# Linear regression is not perfect



$$\hat{y}(x) = p_0 + p_1x + \cdots + p_mx^m$$

Model might perform worse  
in testing in the presence of  
outliers



# Improving Linear Regression

$$W^* = \arg \min_W L(X, W)$$


$$J_2 = \frac{1}{n} \sum_{i=1}^n (f(\tilde{x}_i, W) - y_i)^2 \quad \text{Mean squared error}$$

$$J_1 = \frac{1}{n} \sum_{i=1}^n |f(\tilde{x}_i, W) - y_i| \quad \text{Mean absolute error}$$

$$J_\infty = \max_{1 \leq i \leq n} |f(\tilde{x}_i, W) - y_i| \quad \text{Max error}$$



# Improving Linear Regression



Minimizes


$$(y - Xp)^T (y - Xp)$$

Solution

$$p = (X^T X)^{-1} X^T y$$

*Unbiased*

*High variance*



$$(y - Xp)^T (y - Xp) + \lambda |p|^2$$

$$p = (X^T X + \lambda I)^{-1} X^T y$$


*Biased*

*Low variance*



# Next episode in EE P 596...


$$y = \sigma(\vec{w}^T \vec{x} + b)$$




Classification

$$\nabla_{\vec{w}} L(\hat{y}, y) = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} \nabla_{\vec{w}} z$$

FWD



BWD



Optimizations in Deep Learning using  
Back-propagation and Gradient descent