



LECTURE 3: CONVOLUTIONAL NEURAL NETWORKS


University of Washington, Seattle

Fall 2025




Previously in EE 596...

$$y = \sigma(\vec{w}^T \vec{x} + b)$$






Previously in EE 596...



$$J = L(\vec{w}, b, y)$$

$$\vec{w}_{k+1} = \vec{w}_k - \alpha \nabla_{\vec{w}} J(\vec{w}_k; b)$$

$$b_{k+1} = b_k - \alpha \frac{\partial}{\partial b} J(\vec{w}; b_k)$$






Previously in EE 596...


Training
Dataset

SGD




1 iteration =
Fwd/bwd pass 1 training sample

Mini-batch



1 iteration =
Fwd/bwd pass n-training samples
($n < \text{total \# of training samples}$)



Batch GD



1 iteration =
Fwd/bwd pass full training samples



Previously in EE 596...





Previously in EE 596...

Optimizers

- Vanilla SGD
- Momentum
- AdaGrad
- RMSProp
- Adam



Optimization Techniques

- Data splitting (Train/Val/Test)
- Regularization
- Data normalization
- Batch-normalization
- Network initialization
- Hyperparameter tunings



OUTLINE

Part 1: Need for CNNs

- Limitation of MLP
- Convolution Layer

Part 2: Convolution Filters

- 2D convolution
- Stride
- Padding
- Volume convolutions

Part 3: Composing Convolutional Neural Networks

- Convolution Layer
- Pooling Layers
- Benefits and challenges of CNNs
- Historical CNN examples

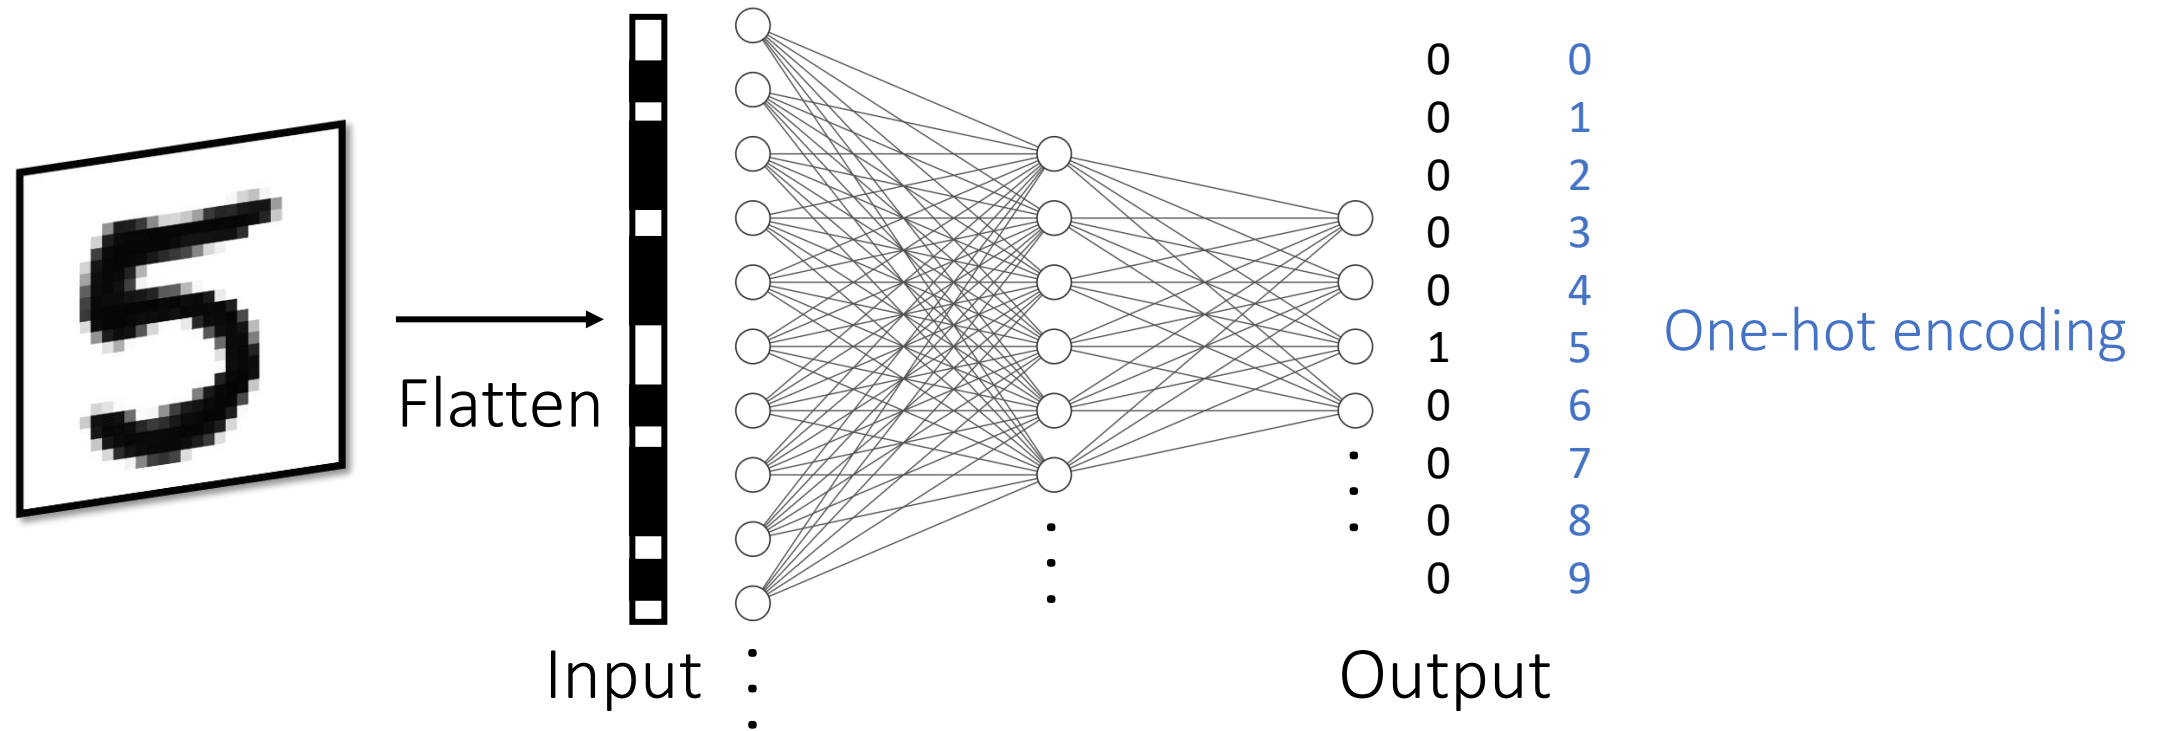


PART 1:

Need for CNNs




MLP for Image Classification (Lab 2)






MLP for Image Classification (Lab 2)

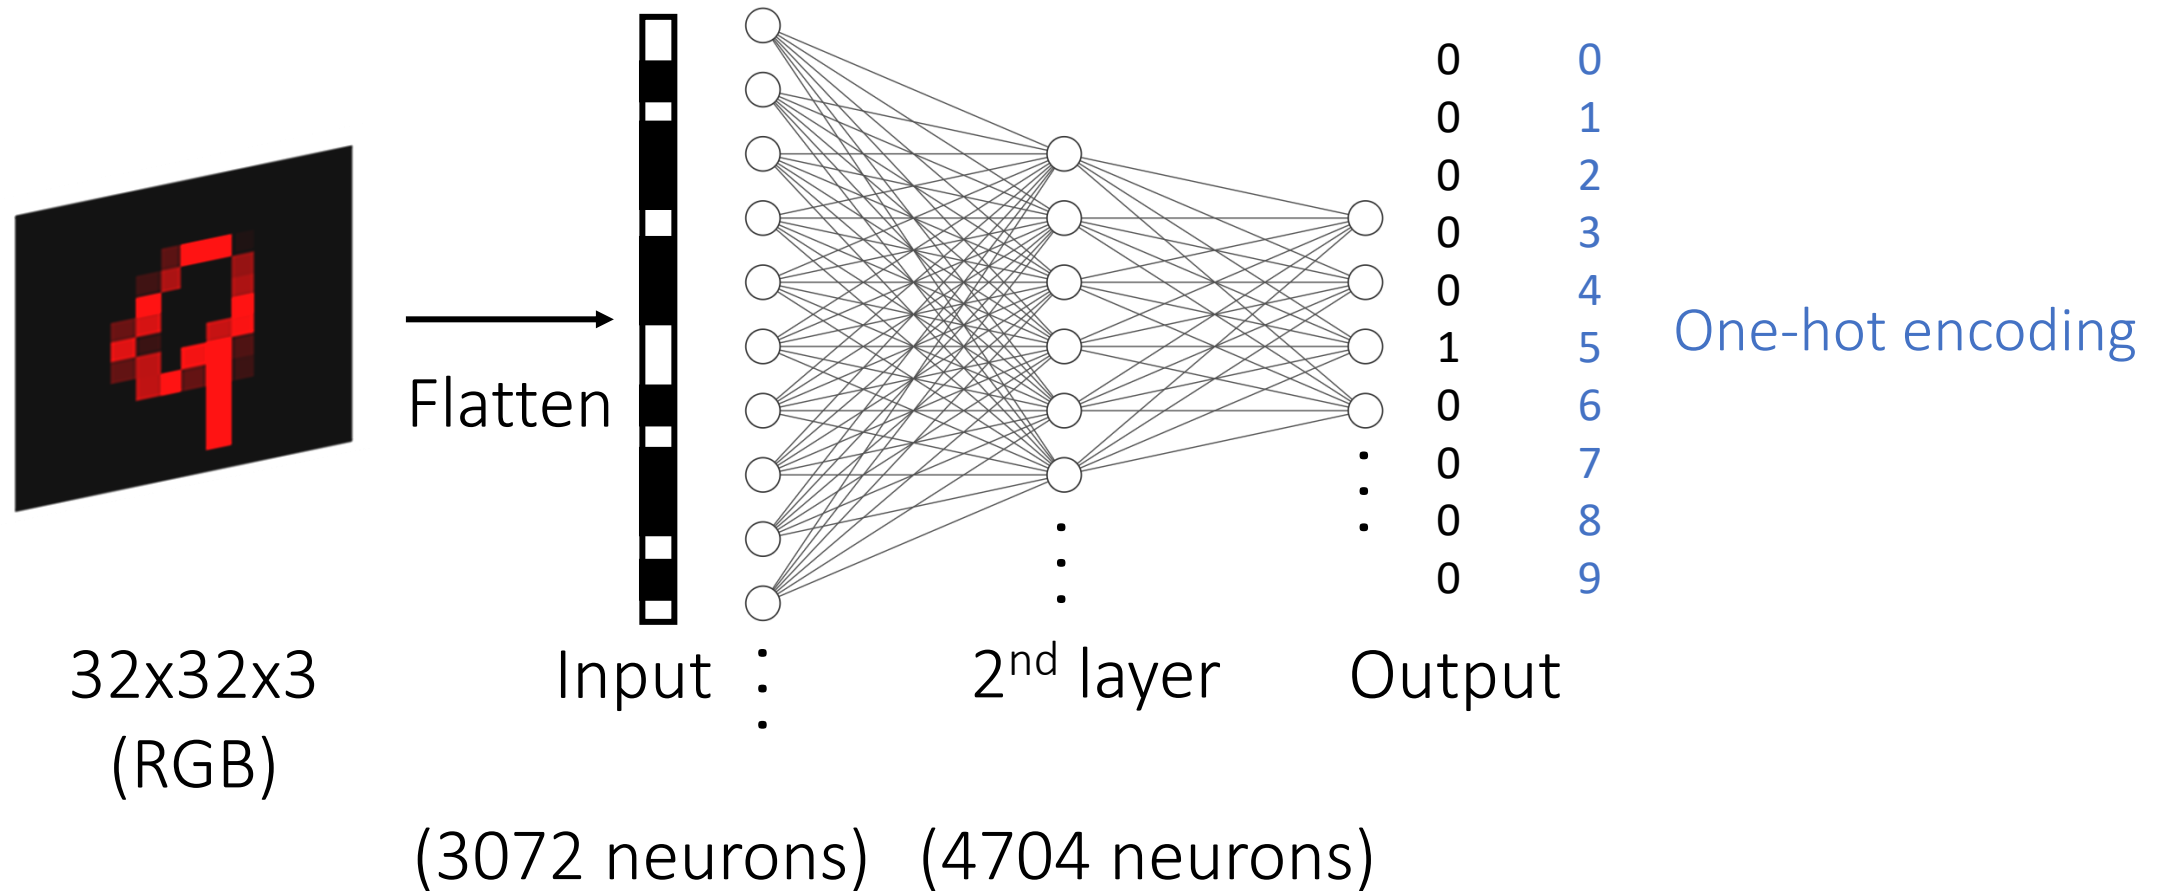





MLP for Image Classification (Lab 2)



MLP for Image Classification (Lab 2)



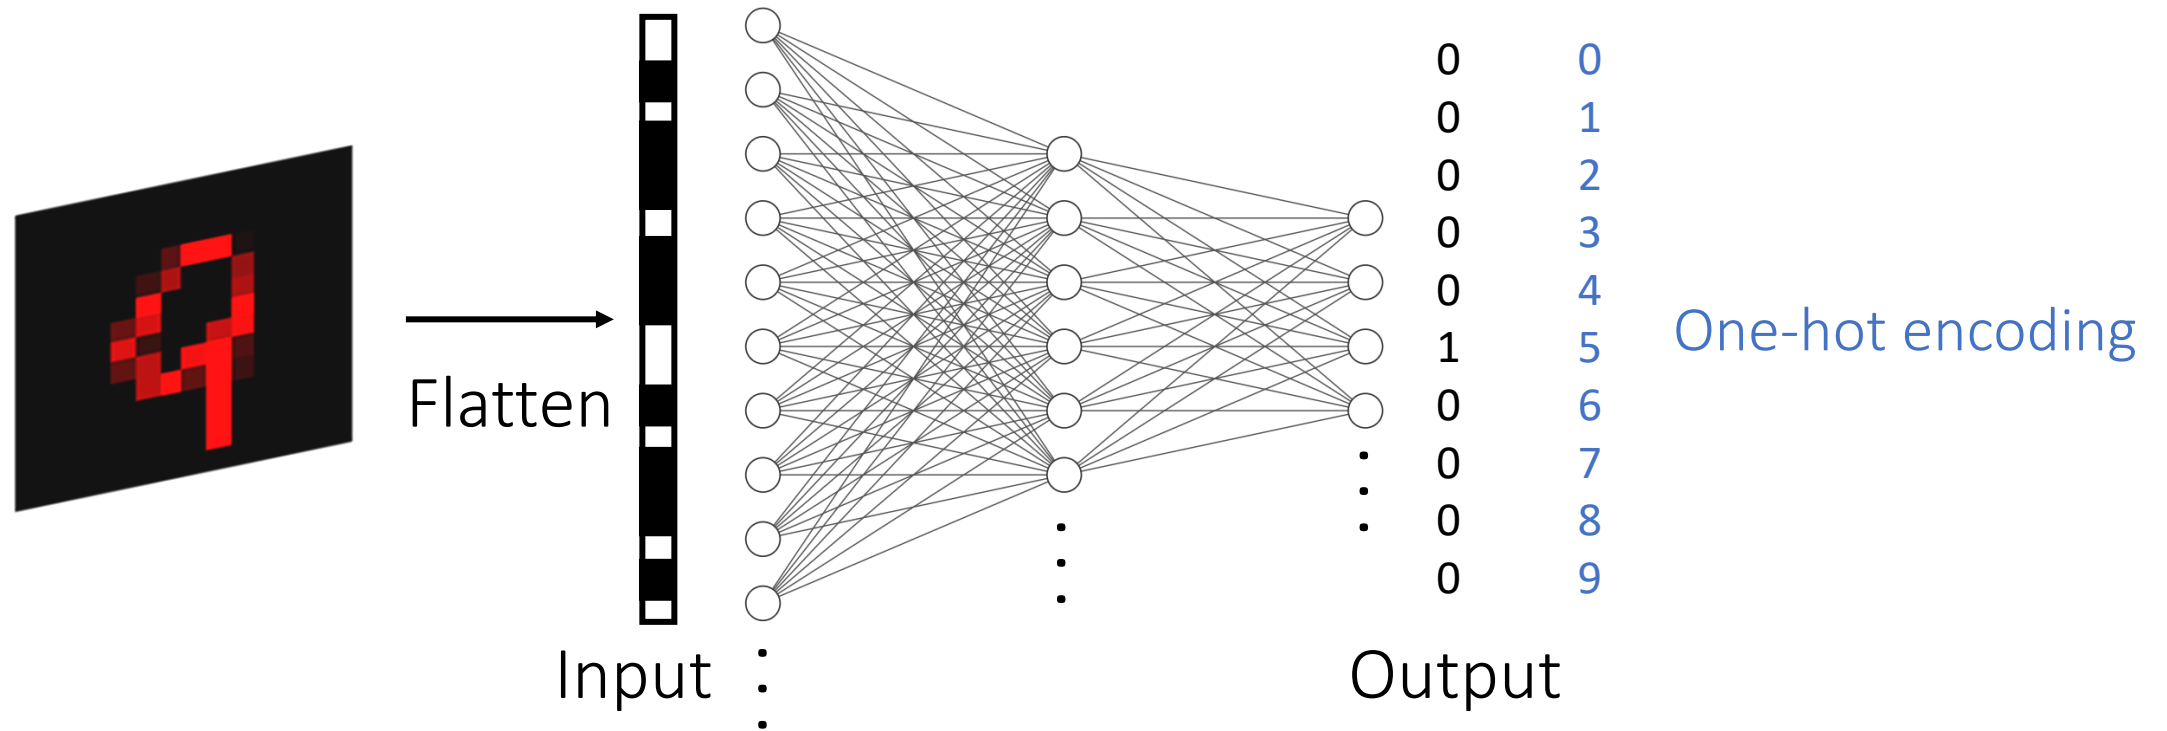
MLP for Image Classification (Lab 2)



(3072 neurons) (4704 neurons)

$$W^{[1]} = [n^{[l-1]}, n^{[l]}] = [3072, 4704] \approx 14M$$

MLP for Image Classification (Lab 2)



Great at Classification


Not as good with Extracting image features

Too many parameters when Flattening images

Specialized Layers for Feature Extractions



Specialized Layers for
Image Feature Extraction



Fully connected layers
(Classifier)



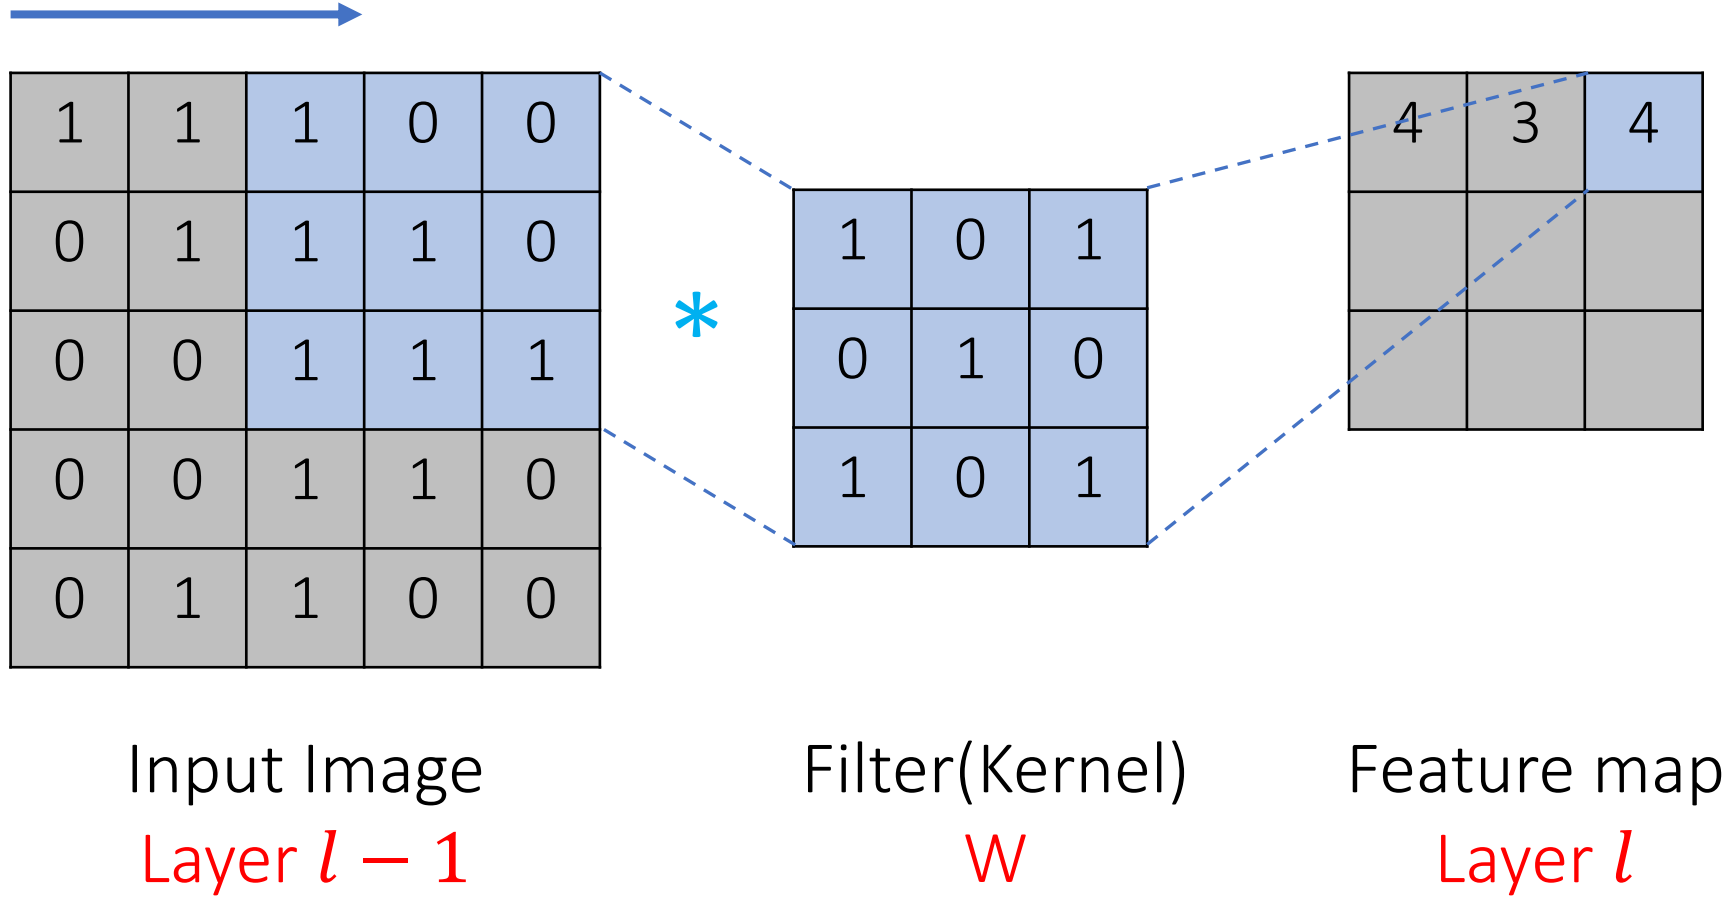


PART 2:

Convolution Filters



Image Convolution



$$\begin{aligned} & (1 * 1) + (0 * 0) + (0 * 1) + \\ & (0 * 1) + (1 * 1) + (0 * 0) + \\ & (1 * 1) + (0 * 0) + (1 * 1) + \end{aligned}$$



Traditional Convolution Filters

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	



Traditional Convolution Filters

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	

CNNs **Learn** these features instead of us guessing



Traditional Convolution Filters

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	

CNNs **Learn** these features instead of us guessing

1000 filters

$3 \times 3 = 9 \times 1000 = 9K$ parameters



Traditional Convolution Filters

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	

CNNs **Learn** these features instead of us guessing

1000 filters

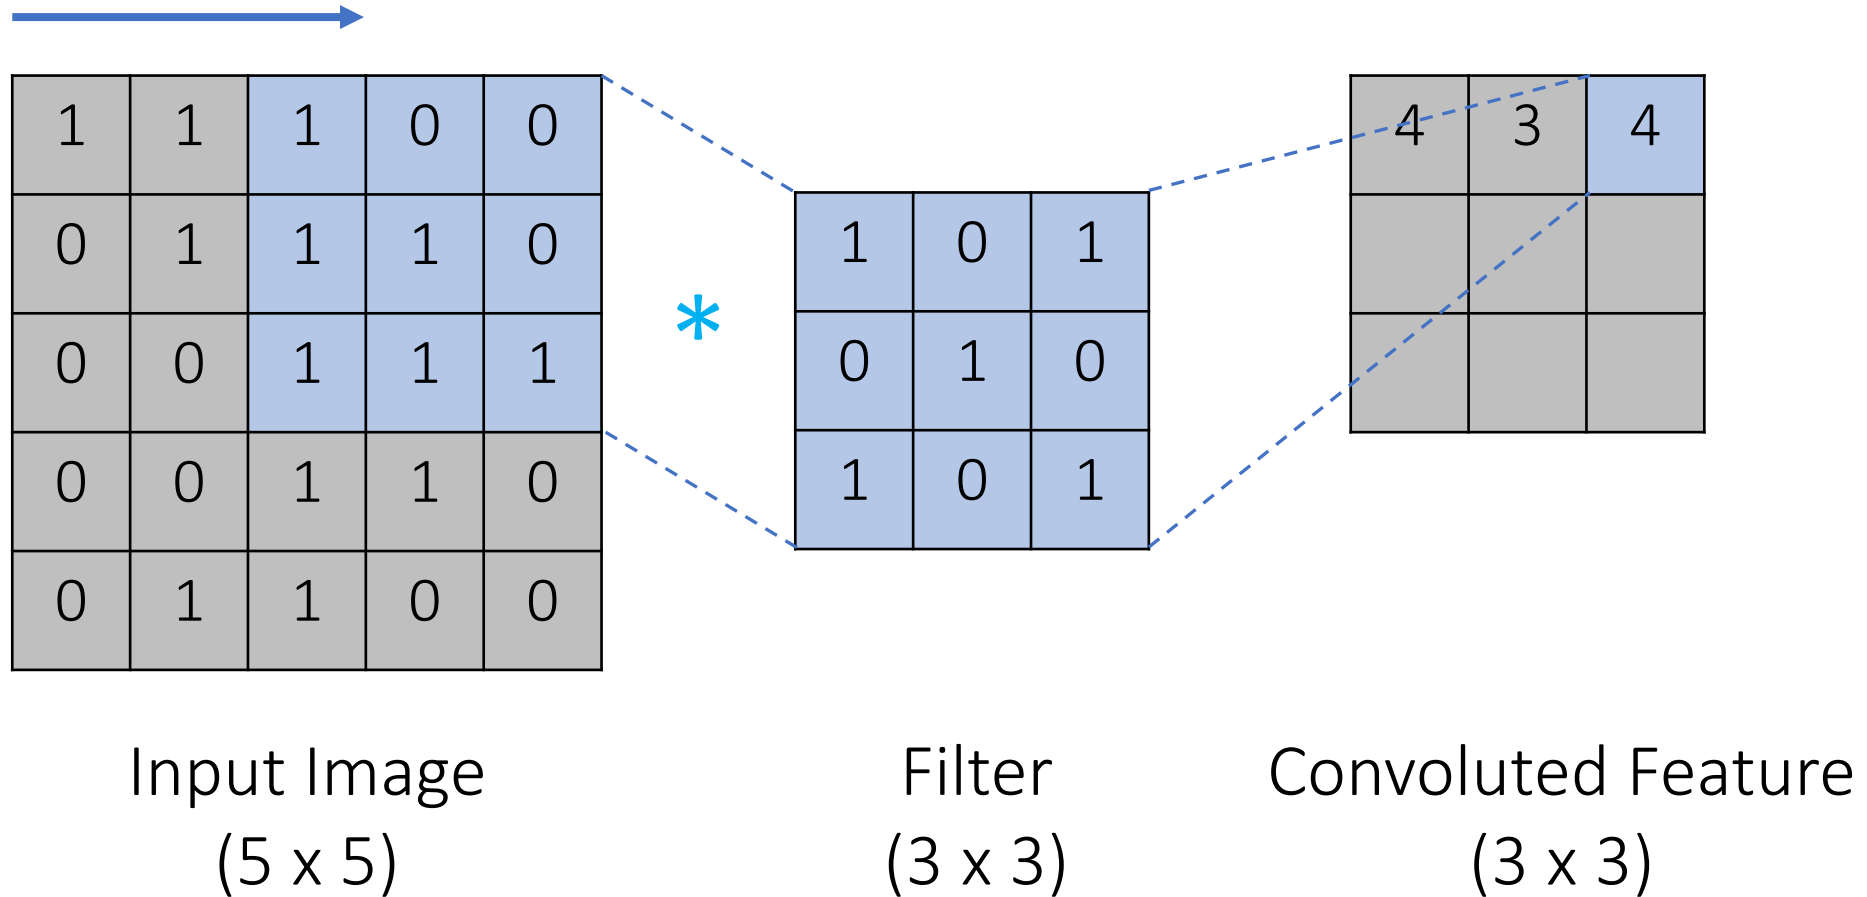
$3 \times 3 = 9 * 1000 = \mathbf{9K}$ parameters

14M vs 9k

Several orders of magnitude of difference in parameters




Convolution Dimensions





Convolution Dimensions





Stride

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

*

1	0	1
0	1	0
1	0	1

Filter


4		

Convolved Feature

Stride = 1



Stride



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

*

1	0	1
0	1	0
1	0	1

Filter


4	3	

Convolved Feature

Stride = 1



Stride



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

*

1	0	1
0	1	0
1	0	1

Filter

4	3	4


Convolved Feature

Stride = 1



Stride

Input = 5 x 5



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

*

1	0	1
0	1	0
1	0	1

Filter

Output = 3 x 3

4	3	4
3		

Convolved Feature

Stride = 1



Stride

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

*

1	0	1
0	1	0
1	0	1

Filter


4	

Convolved Feature

Stride = 2



Stride



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

*

1	0	1
0	1	0
1	0	1

Filter

4	4


Convolved Feature

Stride = 2



Stride

Input = 5 x 5



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

*

1	0	1
0	1	0
1	0	1

Filter

Output = 2 x 2

4	4
2	

Convolved Feature

Stride = 2



Padding

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image
(5x5)



0	0	0	0	0	0	0
0	1	1	1	0	0	0
0	0	1	1	1	0	0
0	0	0	1	1	1	0
0	0	0	1	1	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0

Padding = 1

Padded Image
(7x7)



Padding

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image
(5x5)



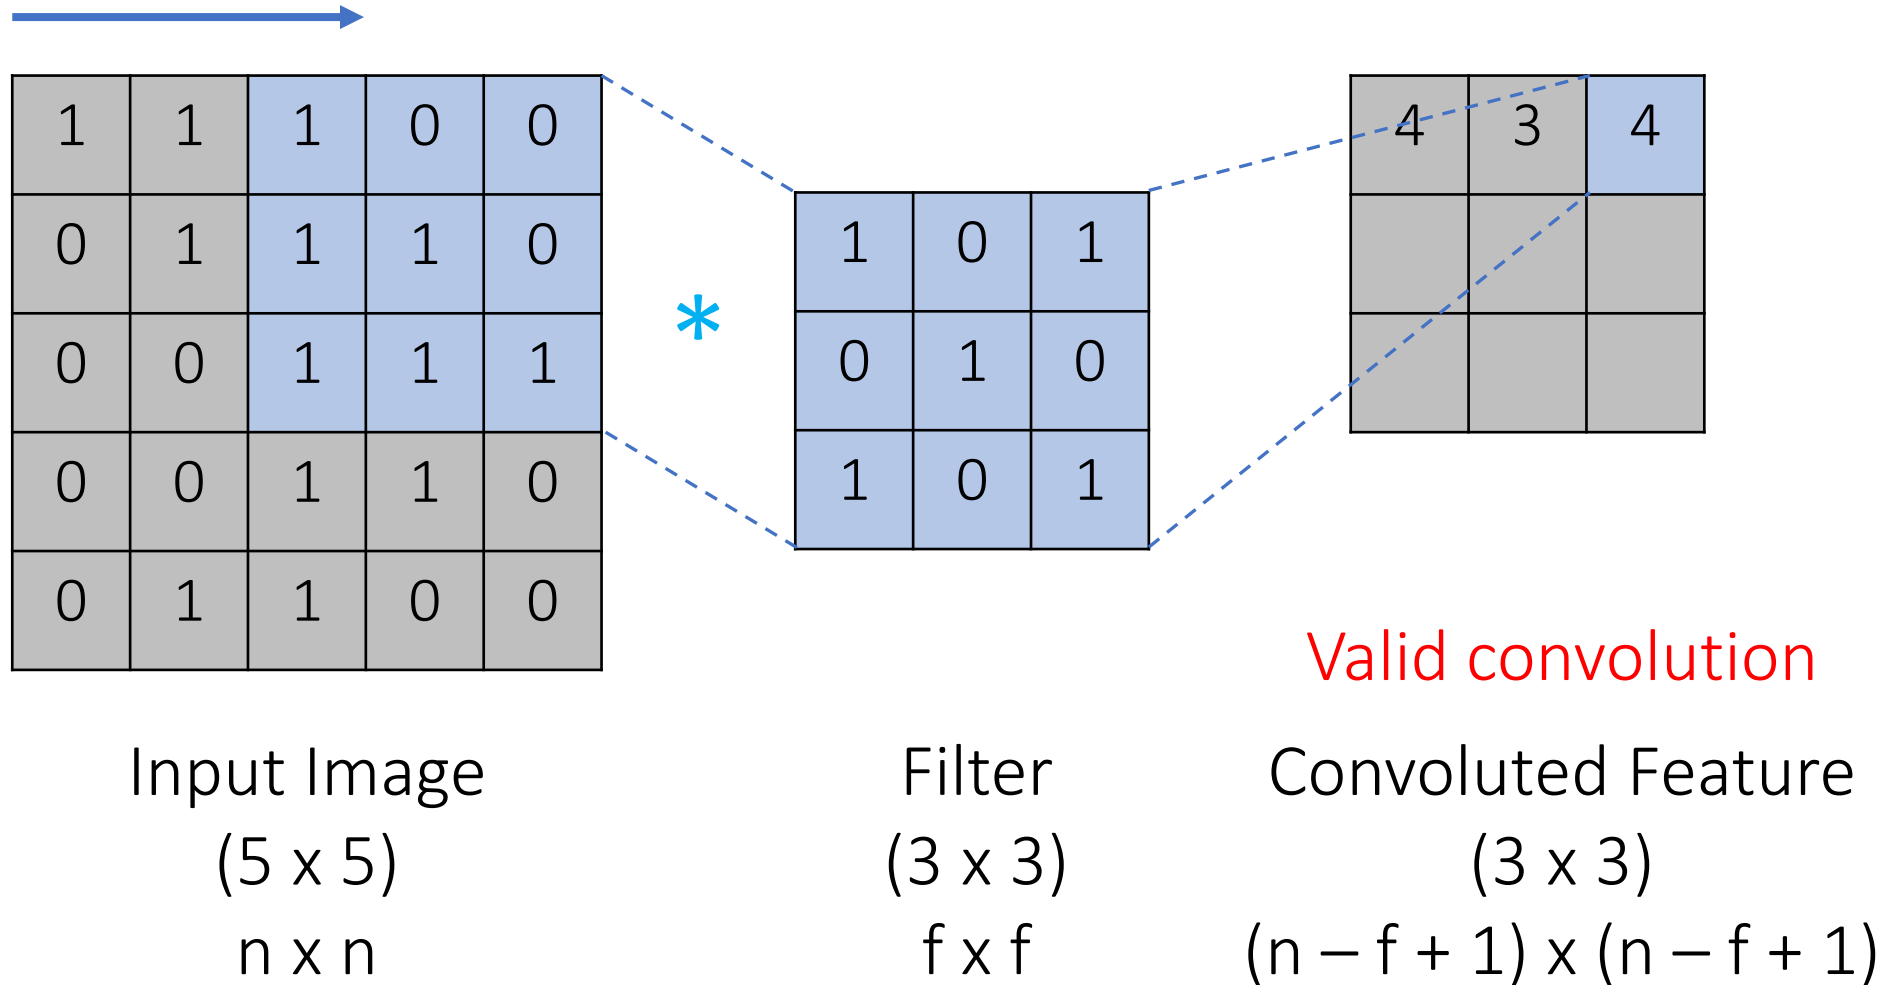
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	1	1	1	0	0	0	0
0	0	0	1	1	1	0	0	0
0	0	0	0	1	1	1	0	0
0	0	0	0	1	1	0	0	0
0	0	0	1	1	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Padding = 2

Padded Image
(9x9)



Convolution Dimensions





Convolution Dimensions



0	0	0	0	0	0	0
0	1	1	1	0	0	0
0	0	1	1	1	0	0
0	0	0	1	1	1	0
0	0	0	1	1	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0

Padded Input Image
(7 x 7)
 $n \times n$



1	0	1
0	1	0
1	0	1

Filter
(3 x 3)
 $f \times f$

2	2	2	1	1
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

Convolved Feature
(5 x 5)
 $(n + 2p - f + 1) \times (n + 2p - f + 1)$



Convolution Dimensions

0	0	0	0	0	0	0
0	1	1	1	0	0	0
0	0	1	1	1	0	0
0	0	0	1	1	1	0
0	0	0	1	1	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0

Padded Input Image
(7 x 7)
 $n \times n$

*

1	0	1
0	1	0
1	0	1

Filter
(3 x 3)
 $f \times f$

2	2	2	1	1
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

Convolved Feature
(5 x 5)

$$(n + 2p - f + 1) \times (n + 2p - f + 1)$$

$$p = (f - 1) / 2$$

Same convolution




Convolution Dimensions

Feature	Valid	Same
Padding	No	Yes
Output size	Smaller than input	Same size as input
Aim	Apply convolution on valid regions only	Preserve spatial dimensions



Generalized Dimensions



$$(n) * (n)$$



$$\left(\frac{n + 2p - f}{s} + 1 \right) * \left(\frac{n + 2p - f}{s} + 1 \right)$$

n : original image dimensions


p : padding size

f : filter dimension

s : stride




Volume Convolution



Input
(5x5x3)



Volume Convolution



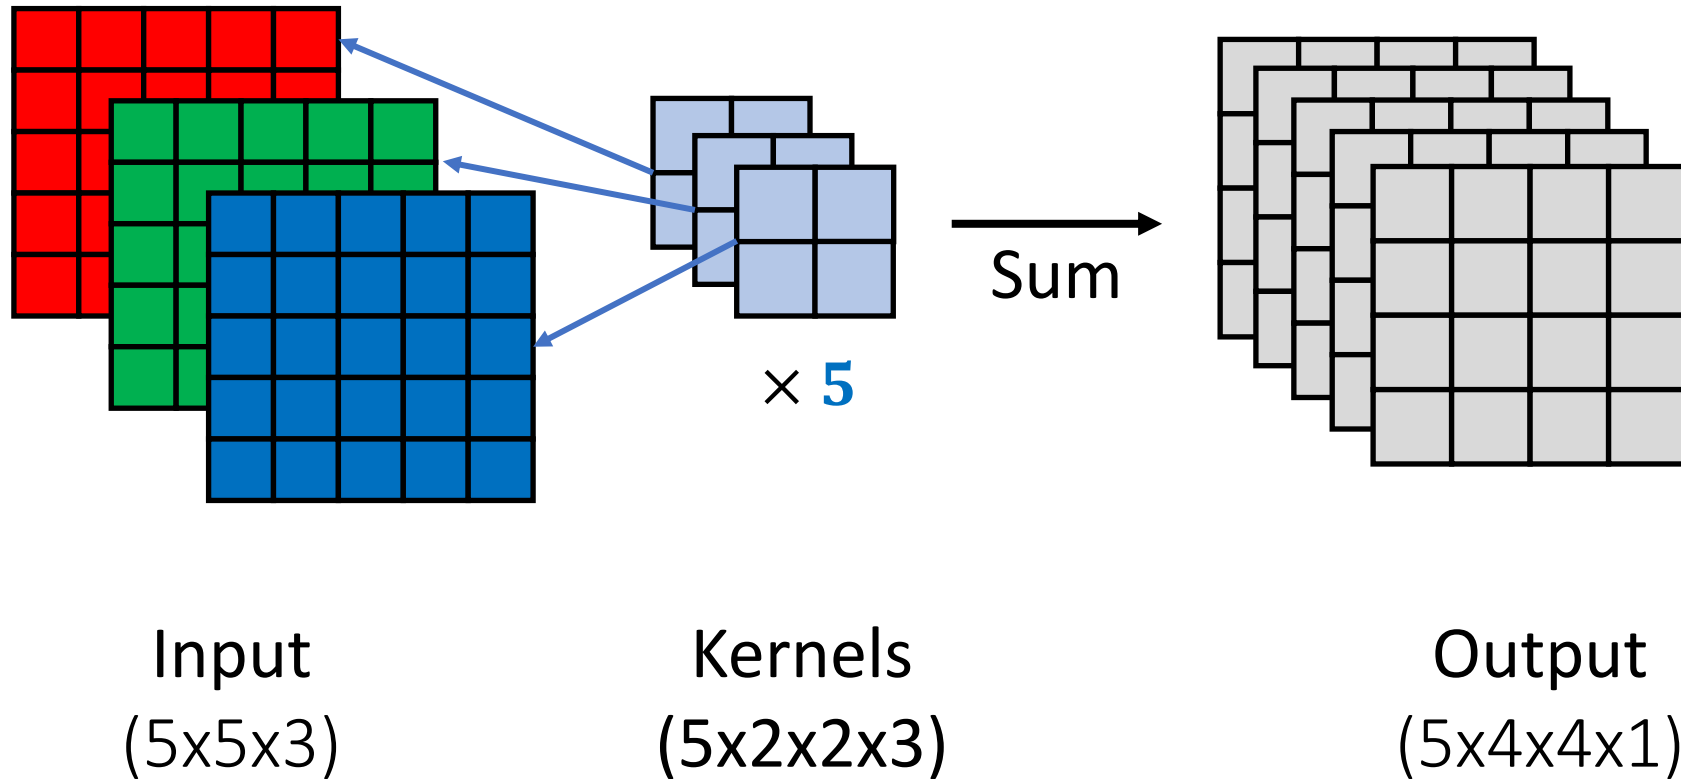
Input
(5x5x3)

Kernel
(2x2x3)

Output
(4x4x1)

(Height x Width x Channels)

Volume Convolution (multiple filters)






PART 3:

Composing CNNs



CNN example







Typical CNN Layers

- **Convolutional Layer (CONV)**
- Pooling Layer (POOL)
- **Fully Connected (FC)**



Typical CNN Layers


- **Convolutional Layer (CONV)**
- Pooling Layer (POOL)
- **Fully Connected (FC)**





Typical CNN Layers

- Convolutional Layer (CONV)
- **Pooling Layer (POOL)**
- Fully Connected (FC)





Max Pooling and Average Pooling

4	9	2	5
5	6	2	4
2	4	5	4
5	6	8	4



9	5
6	8

Max pool

4	9	2	5
5	6	2	4
2	4	5	4
5	6	8	4



6.0	3.3
4.3	5.3

Average pool



Pooling Layers

1	2	5	10	7
0	5	1	4	0
0	4	9	3	15
0	0	2	4	5
0	7	2	0	0

Input Image

Max pool

Dim= 3 x 3
Stride = 1

9		

Pooled Feature



Pooling Layers

1	2	5	10	7
0	5	1	4	0
0	4	9	3	15
0	0	2	4	5
0	7	2	0	0

Input Image

Max pool

Dim= 3 x 3
Stride = 1

9	10	

Pooled Feature



Pooling Layers

1	2	5	10	7
0	5	1	4	0
0	4	9	3	15
0	0	2	4	5
0	7	2	0	0

Input Image

Max pool

Dim= 3 x 3
Stride = 1

9	10	15

Pooled Feature



Pooling Layers

1	2	5	10	7
0	5	1	4	0
0	4	9	3	15
0	0	2	4	5
0	7	2	0	0

Input Image

Avg pool

Dim= 3 x 3
Stride = 1

3		

Pooled Feature



Pooling Layers

1	2	5	10	7
0	5	1	4	0
0	4	9	3	15
0	0	2	4	5
0	7	2	0	0

Input Image

Avg pool

Dim= 3 x 3
Stride = 1

3	4.8	

Pooled Feature



Pooling Layers

1	2	5	10	7
0	5	1	4	0
0	4	9	3	15
0	0	2	4	5
0	7	2	0	0

Input Image

Avg pool

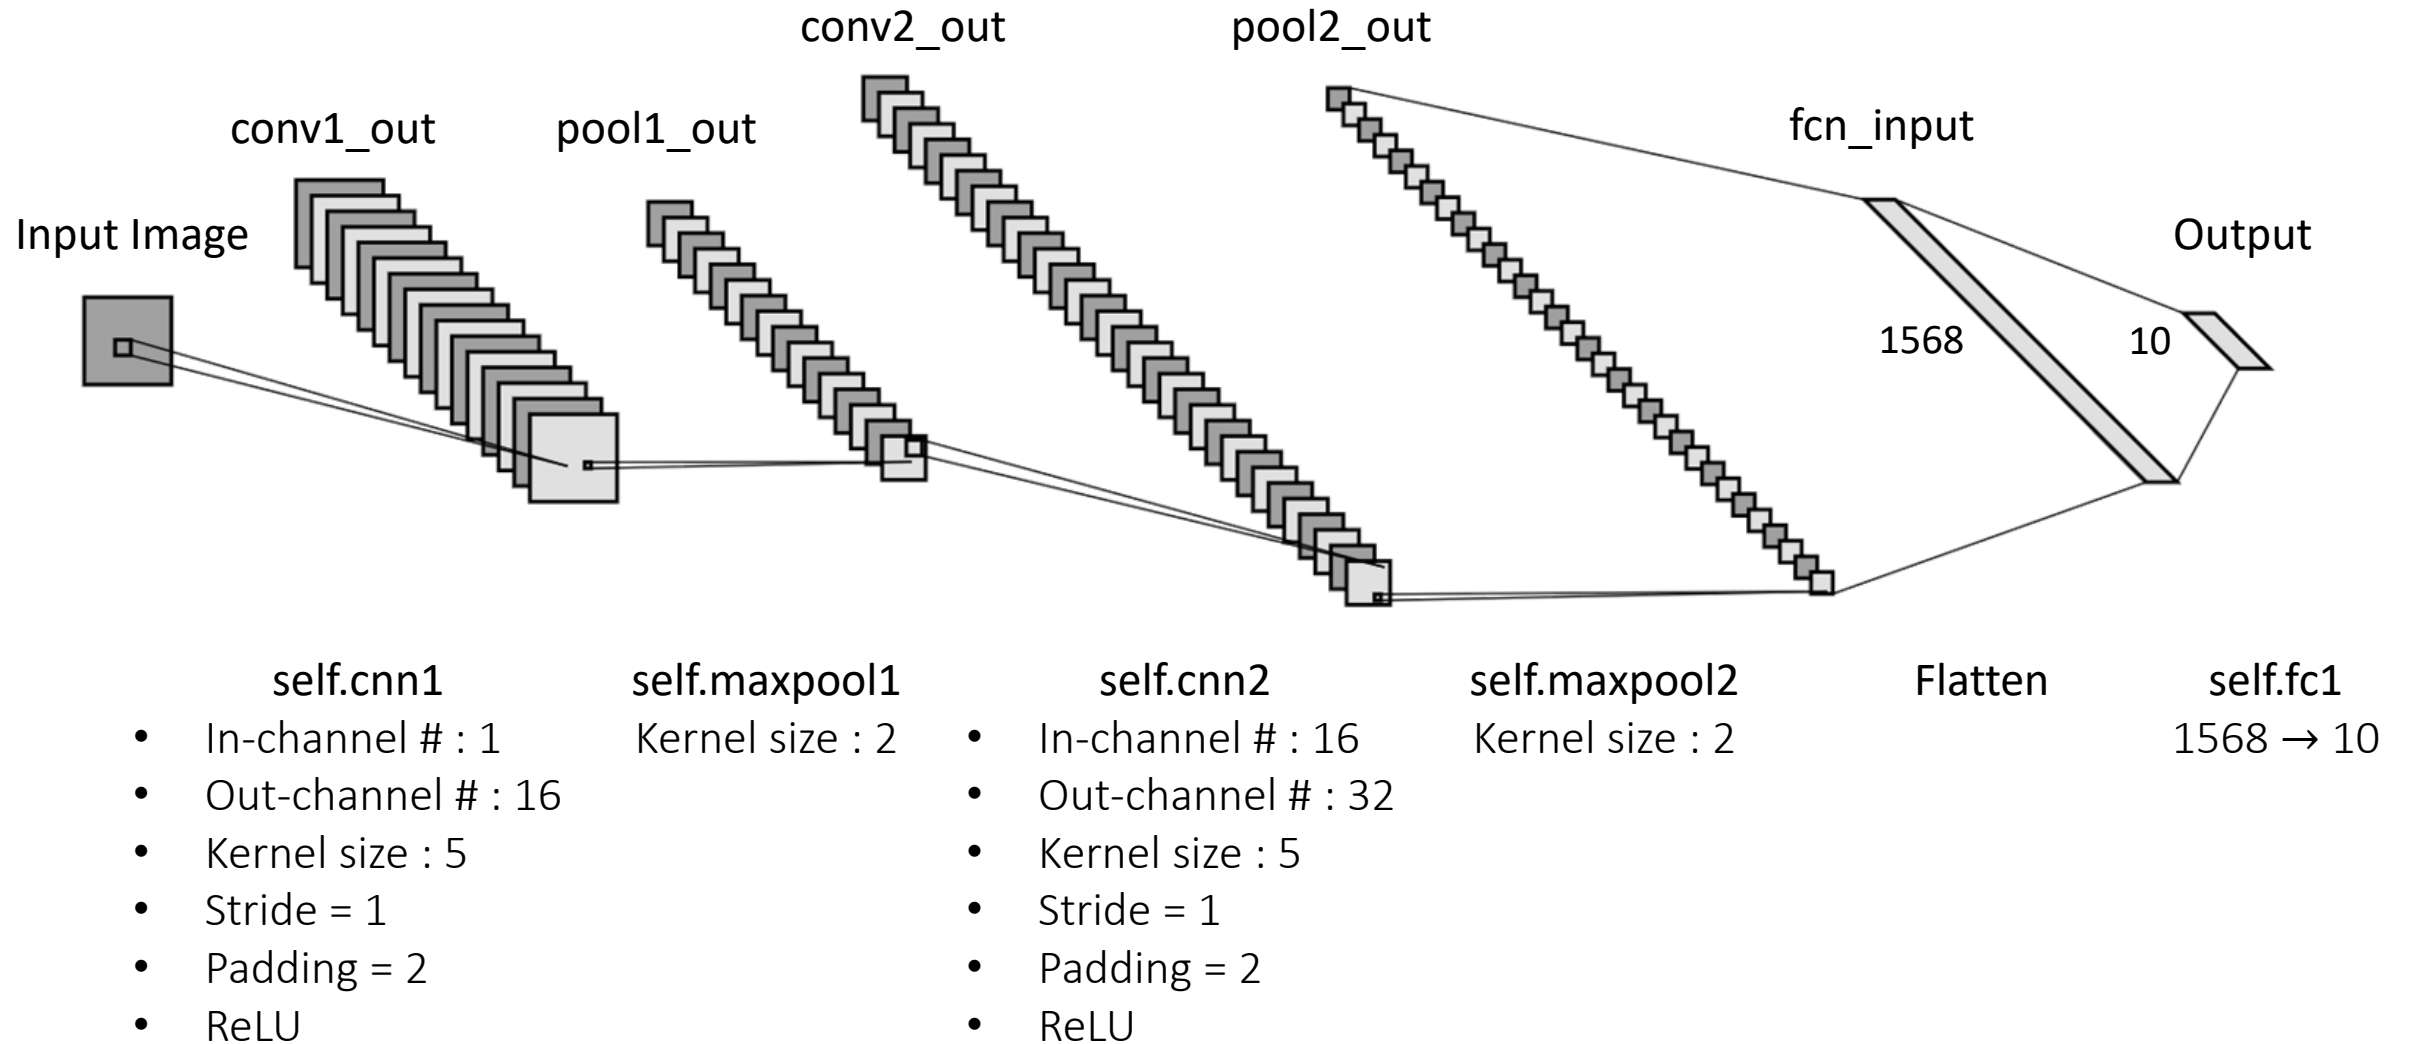
Dim= 3 x 3
Stride = 1

3	4.8	6

Pooled Feature



Full CNN example





Benefits of CNNs

Parameter Sharing

Filter can be useful in different parts of the input (image)

Sparsity of Connections

- In each layer each output value depends only on small number of inputs (Pixel at layer l only depend on subset of pixels in layer $l - 1$)
- Translation invariance (network recognizes patterns regardless of its position in the image)




Benefits of CNNs

Sparsity of Connections

- In each layer each output value depends only on small number of inputs (Pixel at layer l only depend on subset of pixels in layer $l - 1$)
- Translation invariance (network recognizes patterns regardless of its position in the image)

Image credit: Comet



Cat



Cat



Challenges of CNNs

Computational Complexity

Convolutions are expensive $O(N^2n^4)$

Deeper Structure Needed

In each layer each output value depends only on small number of inputs (local)



Popular CNN Architectures (LeNet 5)



Yann LeCun




Leon Bottou



Yoshua Bengio




Patrick Haffner



Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 1998.




LeNet-5 (1998)





LeNet-5 (1998)



Layer 1:

- Convolutional Layer with 6 kernels
- kernel size of 5x5
- Padding = 2, stride = 1

Layer 2:

- Average pooling (2x2 kernel)

Layer 3:

- Convolutional layer with 16 kernels
- kernel size of 5x5
- Padding = 0, stride = 1



LeNet-5 (1998)

Layer 4:

Average pooling (2x2 kernel)

Layer 5:


- Convolutional layer with 120 kernels
- Kernel size of 5x5
- Padding = 0, stride = 1

Layer 6:

- Fully Connected Layer
- Input dimension = 120
- Output dimension = 84

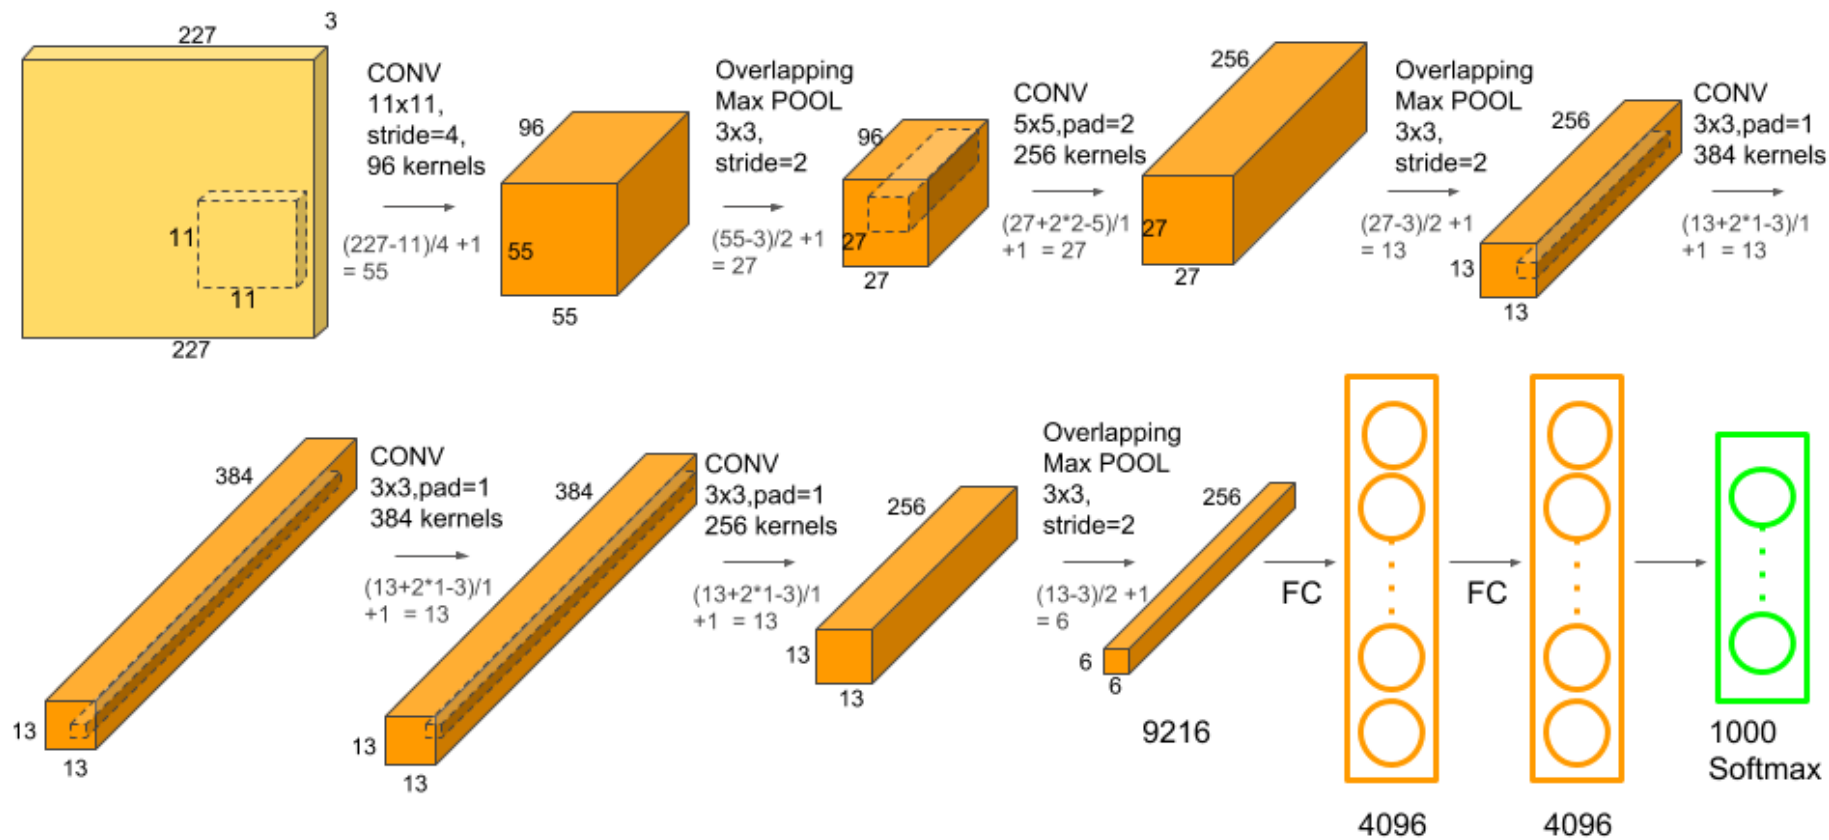
Layer 7:

- Fully Connected Layer
- Input dimension = 84
- Output dimension = 10





AlexNet (2012)






AlexNet (2012)



Alex Krizhevsky



Ilya Sutskever

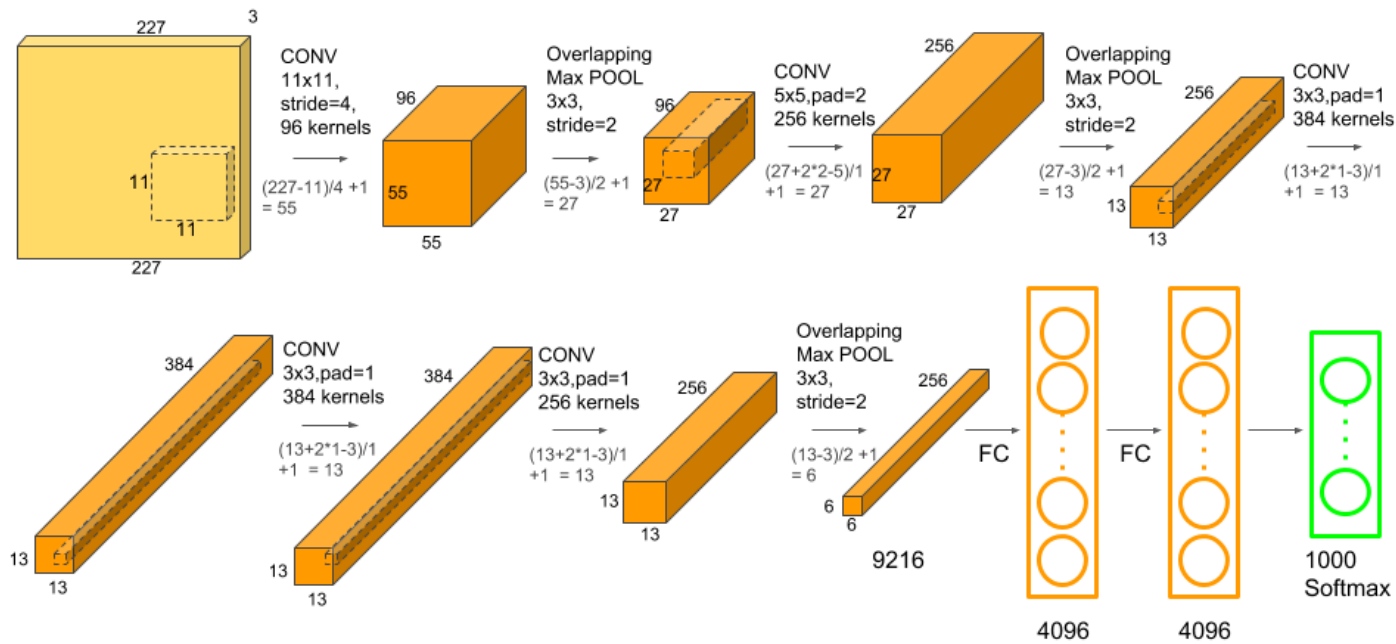


Geoffrey Hinton

Krizhevsky et al., Imagenet classification with deep convolutional neural networks, 2012



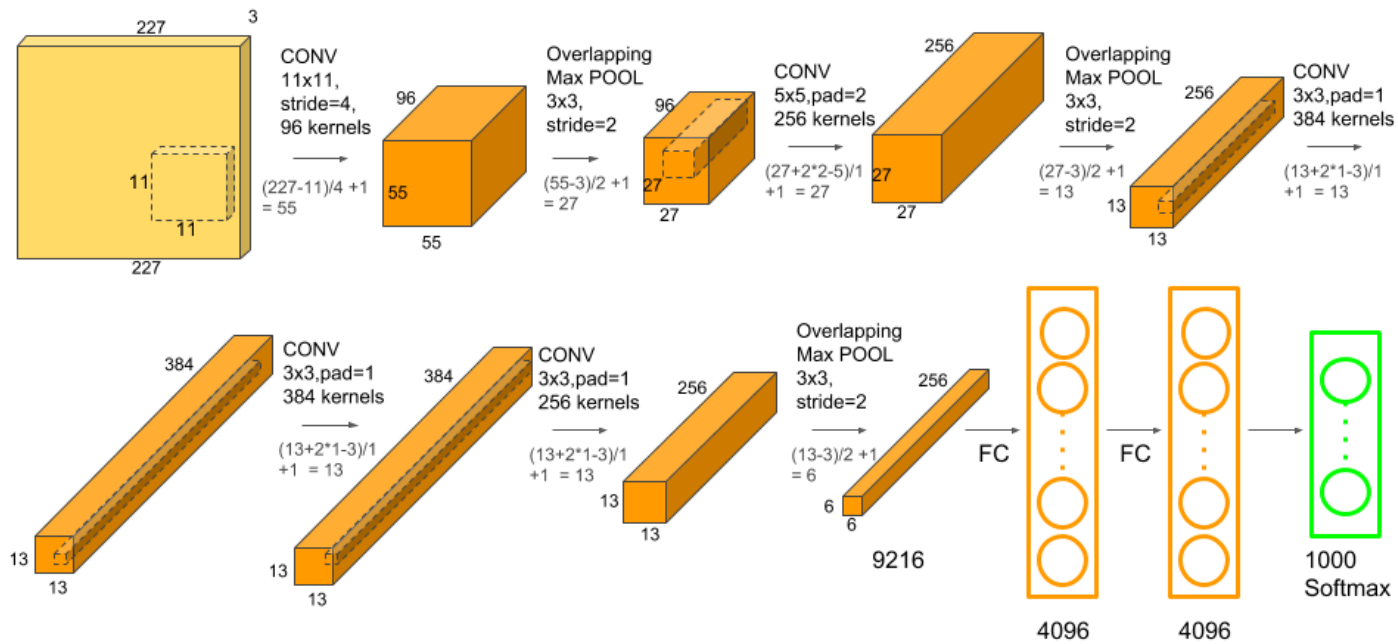
Parameters (AlexNet)



Layer Name	Tensor Size	Weights	Biases	Parameters
Input Image	227x227x3	0	0	0
Conv-1	55x55x96	34,848	96	34,944
MaxPool-1	27x27x96	0	0	0
Conv-2	27x27x256	614,400	256	614,656
MaxPool-2	13x13x256	0	0	0
Conv-3	13x13x384	884,736	384	885,120
Conv-4	13x13x384	1,327,104	384	1,327,488
Conv-5	13x13x256	884,736	256	884,992
MaxPool-3	6x6x256	0	0	0
FC-1	4096x1	37,748,736	4,096	37,752,832
FC-2	4096x1	16,777,216	4,096	16,781,312
FC-3	1000x1	4,096,000	1,000	4,097,000
Output	1000x1	0	0	0
Total				62,378,344



Parameters (AlexNet)



- Much bigger than LeNet (60M parameters)
- ReLU
- Multiple GPUs
- Local Response Normalization (LRN)



VGG-16 (2014)

CONV: f=3, s=1, same

POOL: f=2, s=2

Order: CCP CCP CCCP CCCP CCCP FFS

Nf: 2^6 2^7 2^8 2^9 2^9

~138 mil parameters





VGG-16 (2014)

CONV: $f=3$, $s=1$, same

POOL: $f=2$, $s=2$

Order: CCP CCP CCCP CCCP CCCP FFS

Nf: 2^6 2^7 2^8 2^9 2^9

~138 mil parameters

- Multiple convolution layers
- Smaller convolution filters
- Modularized architecture (VGG-19)

