



# LECTURE 4:

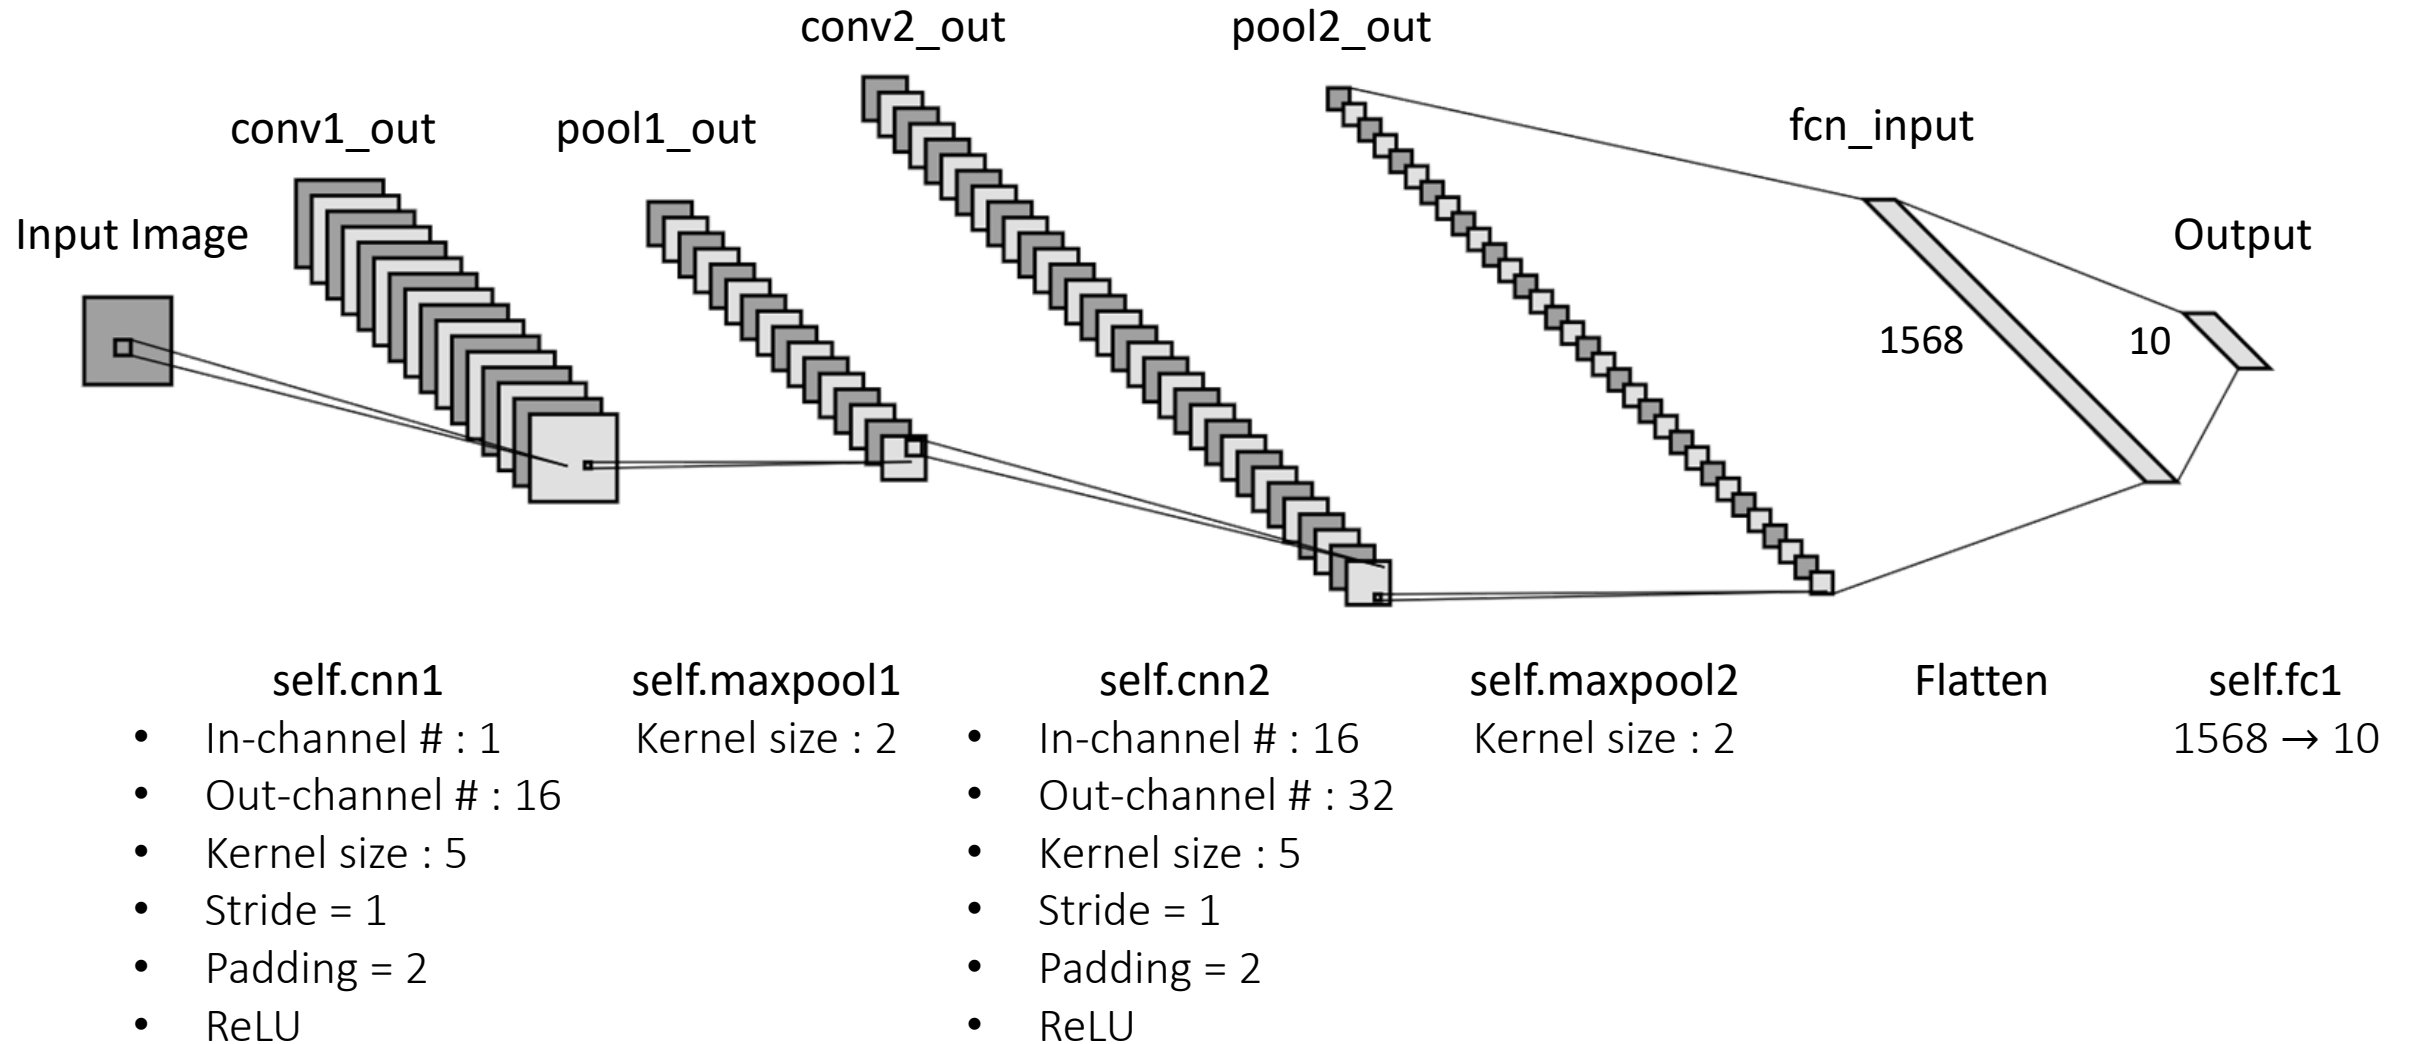
# RECURRENT NEURAL NETWORK

University of Washington, Seattle

Fall 2025



# Previously in EEP 596...





# OUTLINE

## Part 1: Introduction to RNNs

- Why do we need RNNs?
- RNN Architecture
- Embedding and Decoder

## Part 2: Training RNNs

- Backpropagation in RNNs
- Vanishing/Exploding Gradient Problem
- Training with Teacher Forcing

## Part 3: RNN Problem Types

- RNN Configurations
- RNN Extensions



# INTRODUCTION TO RNNs


Why do we need RNNs?

RNN Architecture

Embedding and Decoder




# Why Do We Need RNNs?



Feed-Forward Network




# Why Do We Need RNNs?



Feed-Forward Network



# Why Do We Need RNNs?



Feed-Forward Network neurons have no memory of past inputs



# Why Do We Need RNNs?

Korean

안녕하세요, 제 이름은 지민이에요



English

Hello, my name is Jimin





# Why Do We Need RNNs?

Korean

안녕하세요,

English


Hello,



# Why Do We Need RNNs?

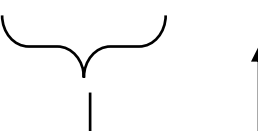
Korean

안녕하세요, 제



English

Hello, my

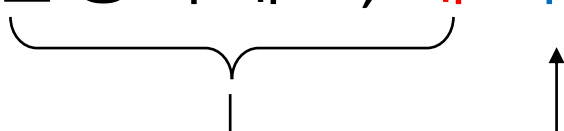




# Why Do We Need RNNs?

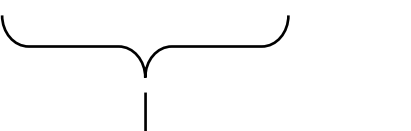
Korean

안녕하세요, 제 이름



English

Hello, my name

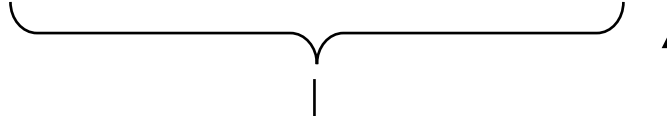




# Why Do We Need RNNs?


Korean

안녕하세요, 제 이름은



English

Hello, my name is





# Why Do We Need RNNs?

Korean

안녕하세요, 제 이름은 지민이에요

The diagram illustrates the sequential nature of Korean. A bracket under '안녕하세요,' is followed by another bracket under '제 이름은', with an arrow pointing from the second bracket to '지민이에요', indicating that the meaning of the second part depends on the first.

English

Hello, my name is Jimin

The diagram illustrates the sequential nature of English. A bracket under 'Hello,' is followed by another bracket under 'my name is', with an arrow pointing from the second bracket to 'Jimin', indicating that the meaning of the second part depends on the first.



# Why Do We Need RNNs?

Korean

안녕하세요, 제 이름은 지민이에요

English

Hello, my name is Jimin

Each word in a sentence is dependent to the past words →

Need memory



# Why Do We Need RNNs?

Korean      안녕하세요, 제 이름은 지민이에요, 그리고 저는 비디오게임을 좋아해요

English      Hello, my name is Jimin, and I like videogames

A sentence (input) could have **different sizes**



# Why Do We Need RNNs?

We need a neural network architecture that can handle:





# Why Do We Need RNNs?

We need a neural network architecture that can handle:

- Data order



# Why Do We Need RNNs?

We need a neural network architecture that can handle:

- Data order
- Temporal dependencies



# Why Do We Need RNNs?

We need a neural network architecture that can handle:


- Data order
- Temporal dependencies
- Variable input sizes



# Why Do We Need RNNs?

We need a neural network architecture that can handle:

- Data order
- Temporal dependencies
- Variable input sizes



Recurrent Neural Network

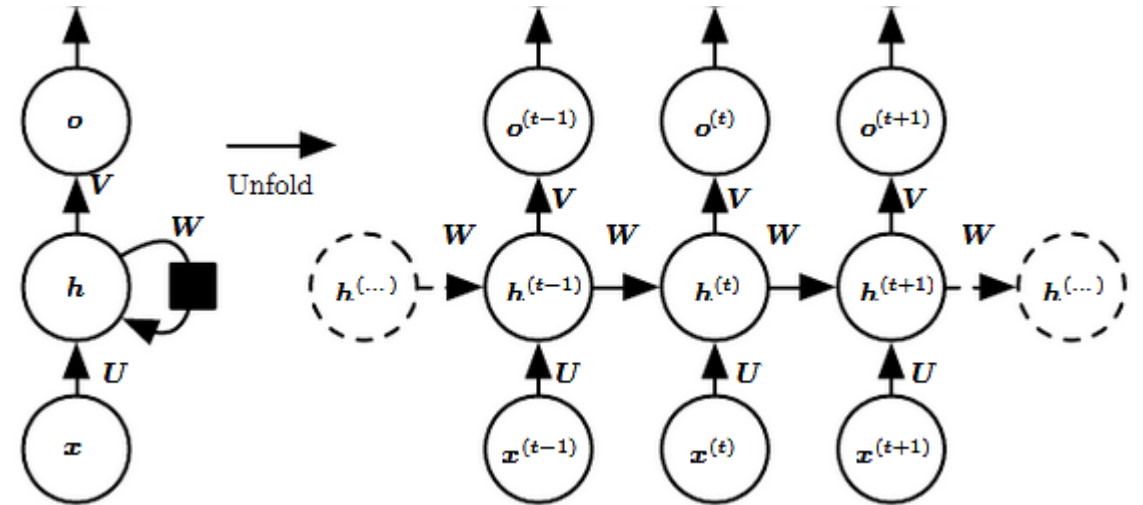



# RNN Architecture

Output

Hidden

Input



Unfold in Time

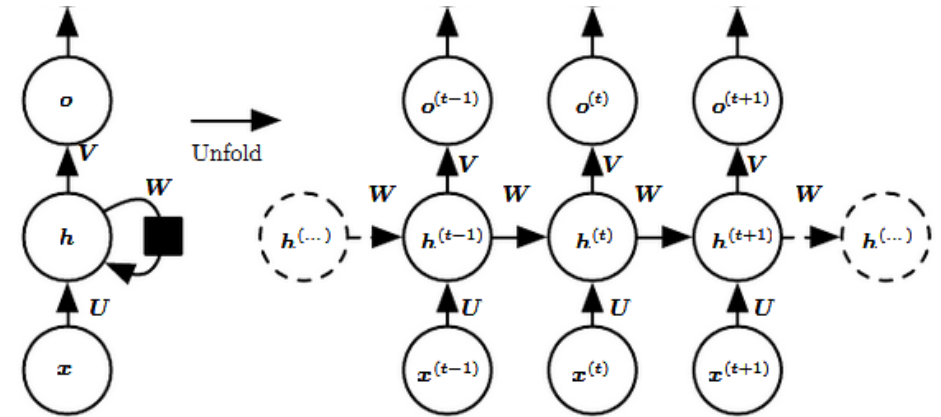



# RNN Architecture

Output

Hidden

Input




Unfold in Time

$$\begin{aligned} a^{(t)} &= b + \mathbf{W}h^{(t-1)} + \mathbf{U}x^{(t)} \\ h^{(t)} &= \tanh(a^{(t)}) \\ o^{(t)} &= c + \mathbf{V}h^{(t)} \\ \hat{y}^{(t)} &= \text{softmax}(o^{(t)}) \end{aligned}$$



# Brain is Highly Recurrent



Neurons themselves have temporal voltage dynamics

Different parts of brain exchange information both forward and backward



# Brain is Highly Recurrent

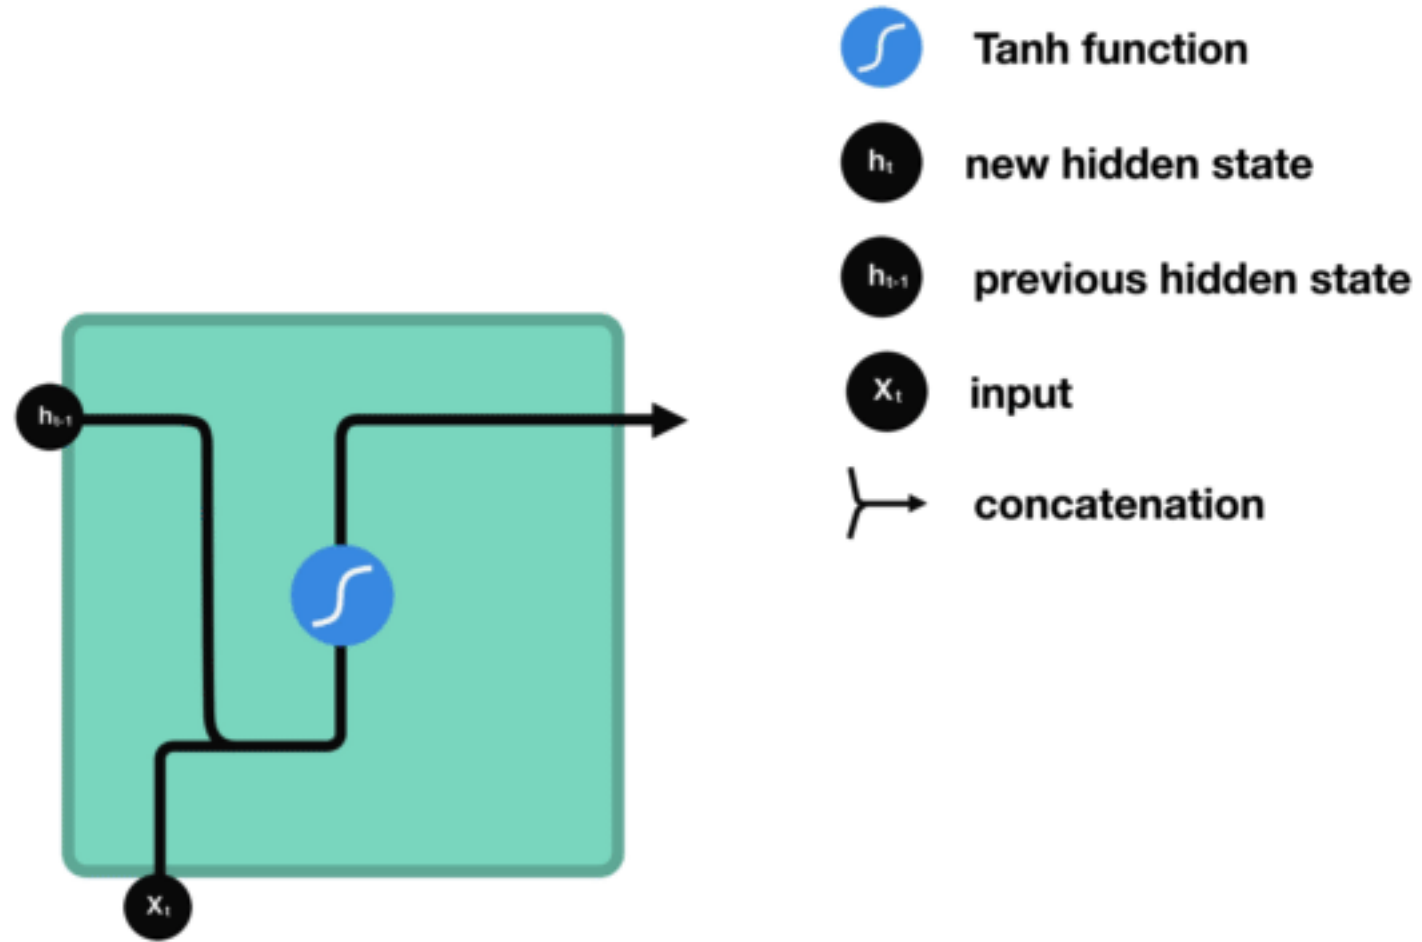


Credit: Allen Institute for Brain Science






# RNN Architecture






# RNN Architecture






# RNN Architecture



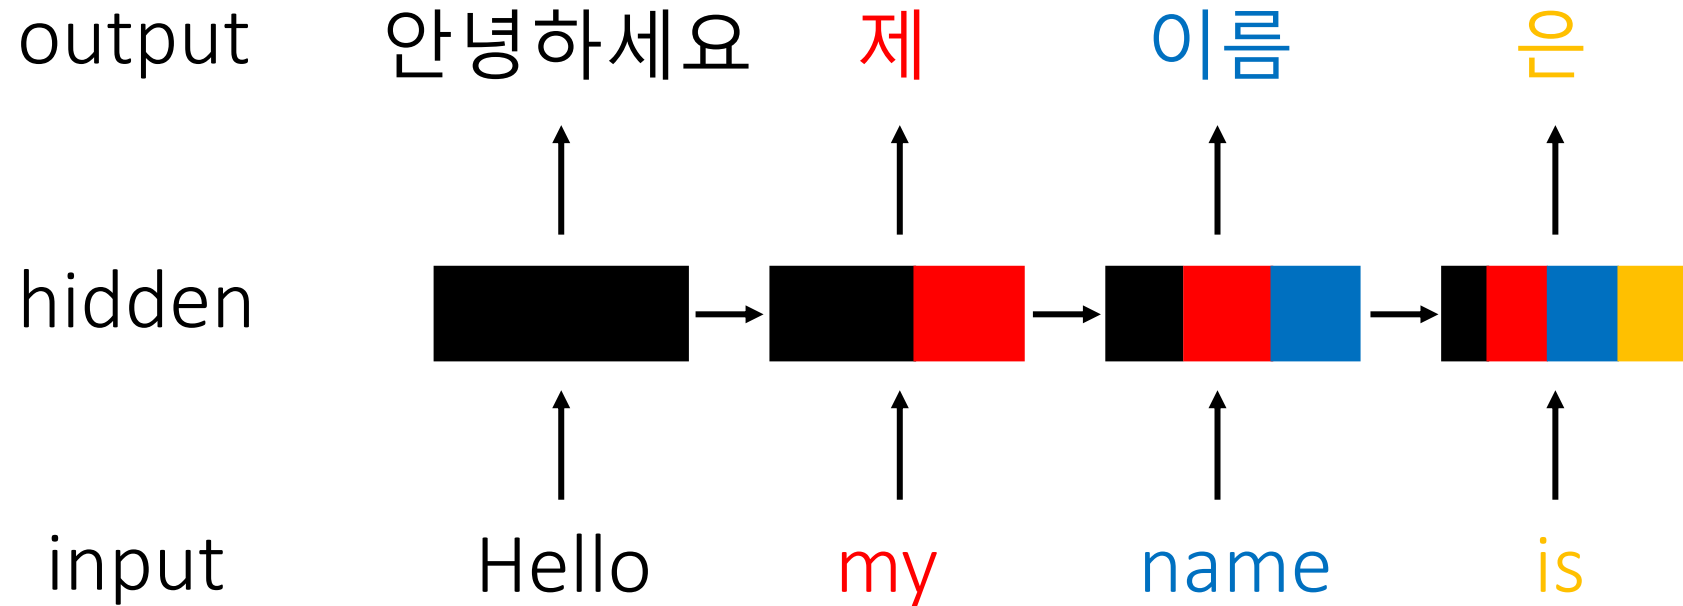


# RNN Architecture



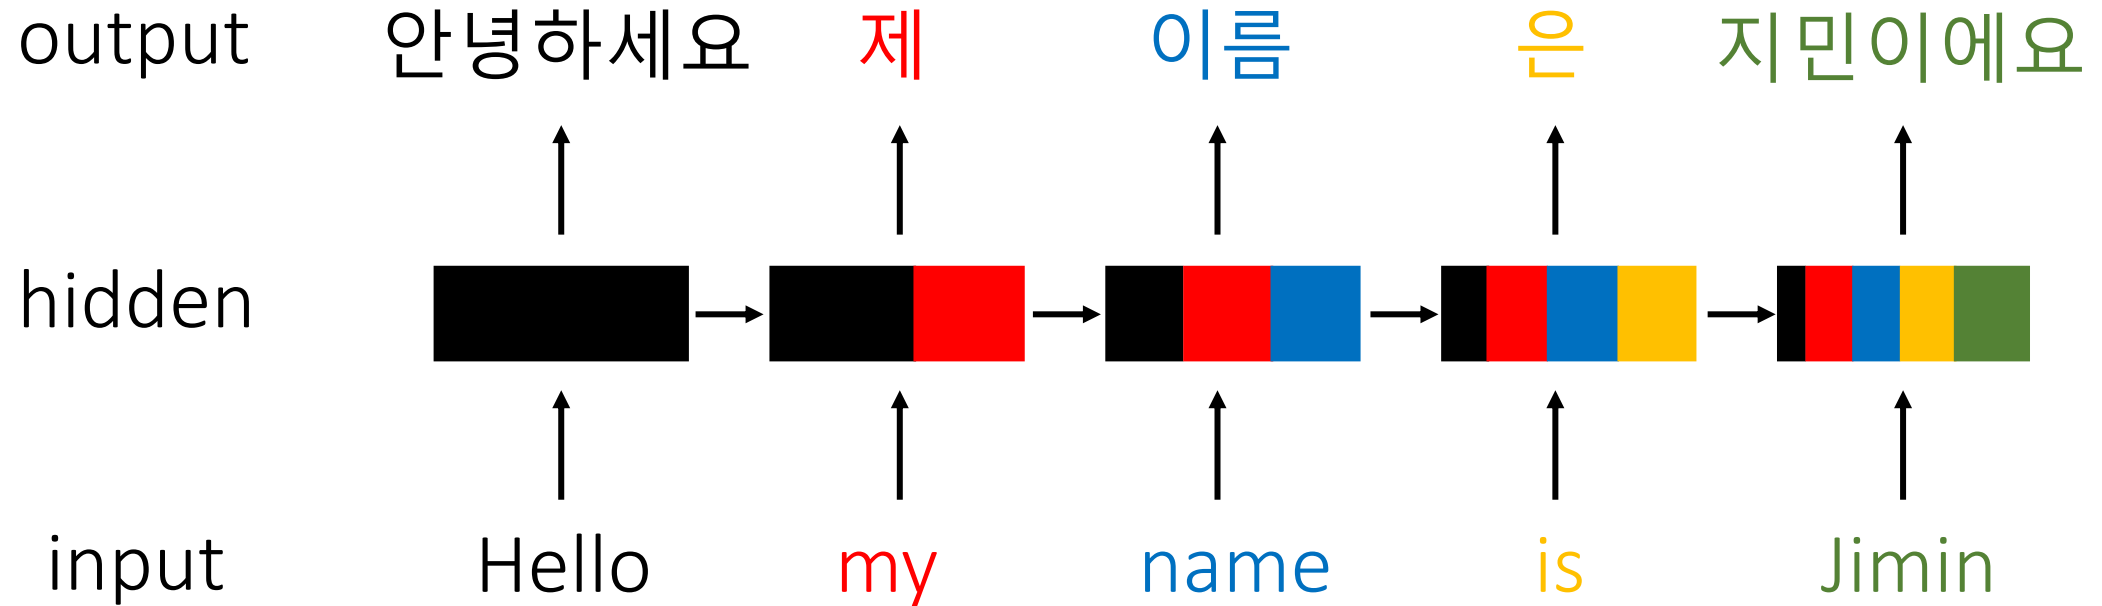


# RNN Architecture





# RNN Architecture





# Sequential Data

Speech recognition



“The quick brown fox jumped  
over the lazy dog.”

Music generation

∅



Sentiment classification

“There is nothing to like  
in this movie.”



DNA sequence analysis

AGCCCCTGTGAGGAACTAG



AG**CCCCTGTGAGGAACT**AG

Machine translation

Voulez-vous chanter avec  
moi?



Do you want to sing with  
me?

Video activity recognition



Running

Name entity recognition


Yesterday, Harry Potter  
met Hermione Granger.



Yesterday, **Harry Potter**  
met **Hermione Granger**.



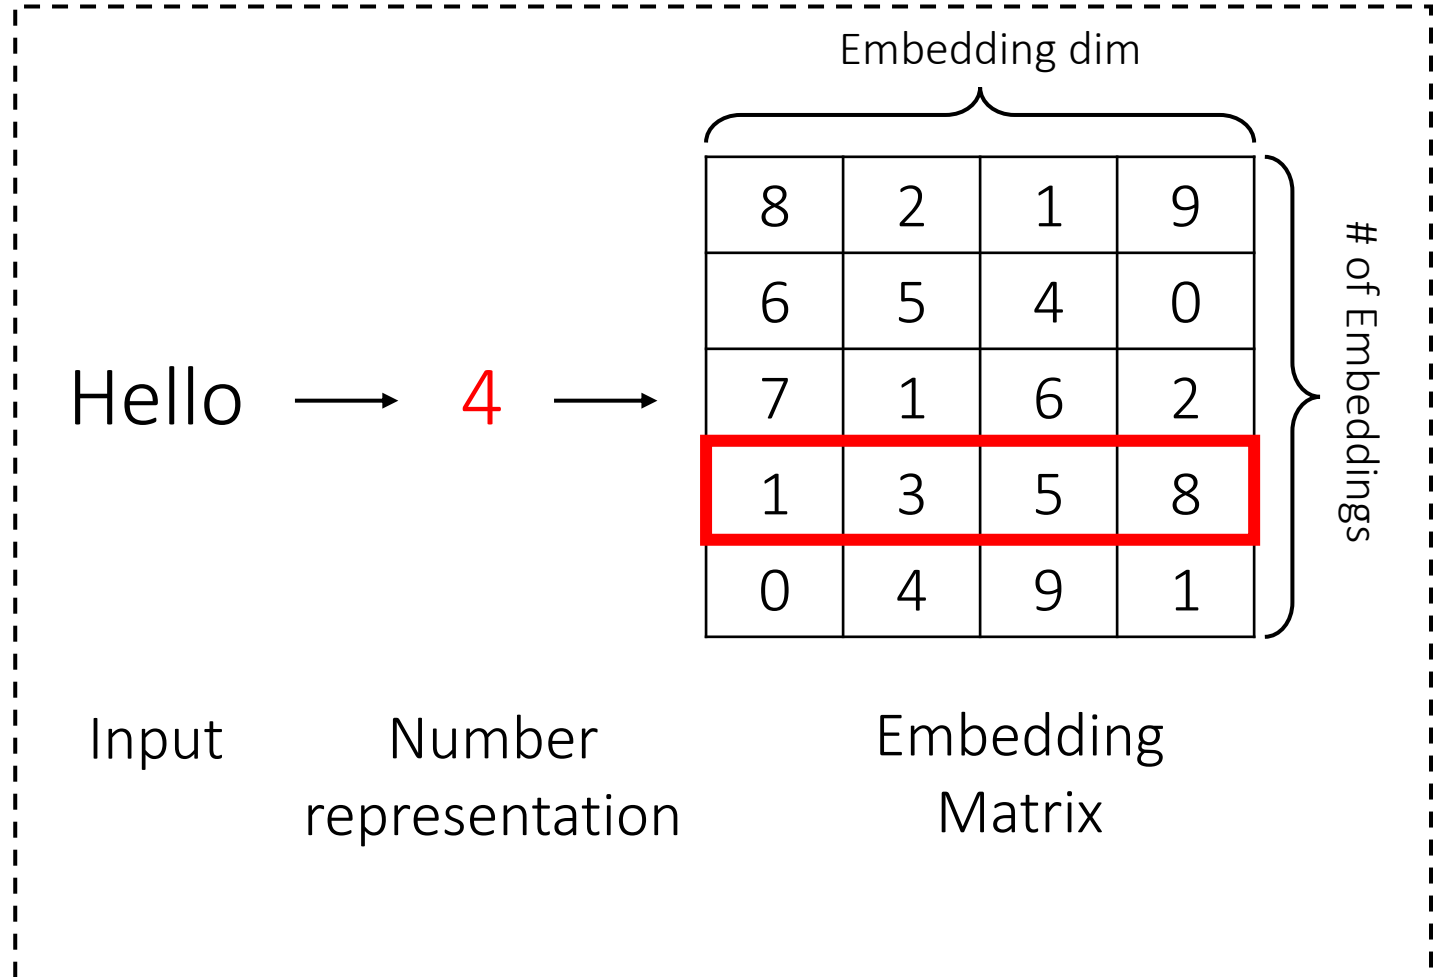

# Embedding and Decoder





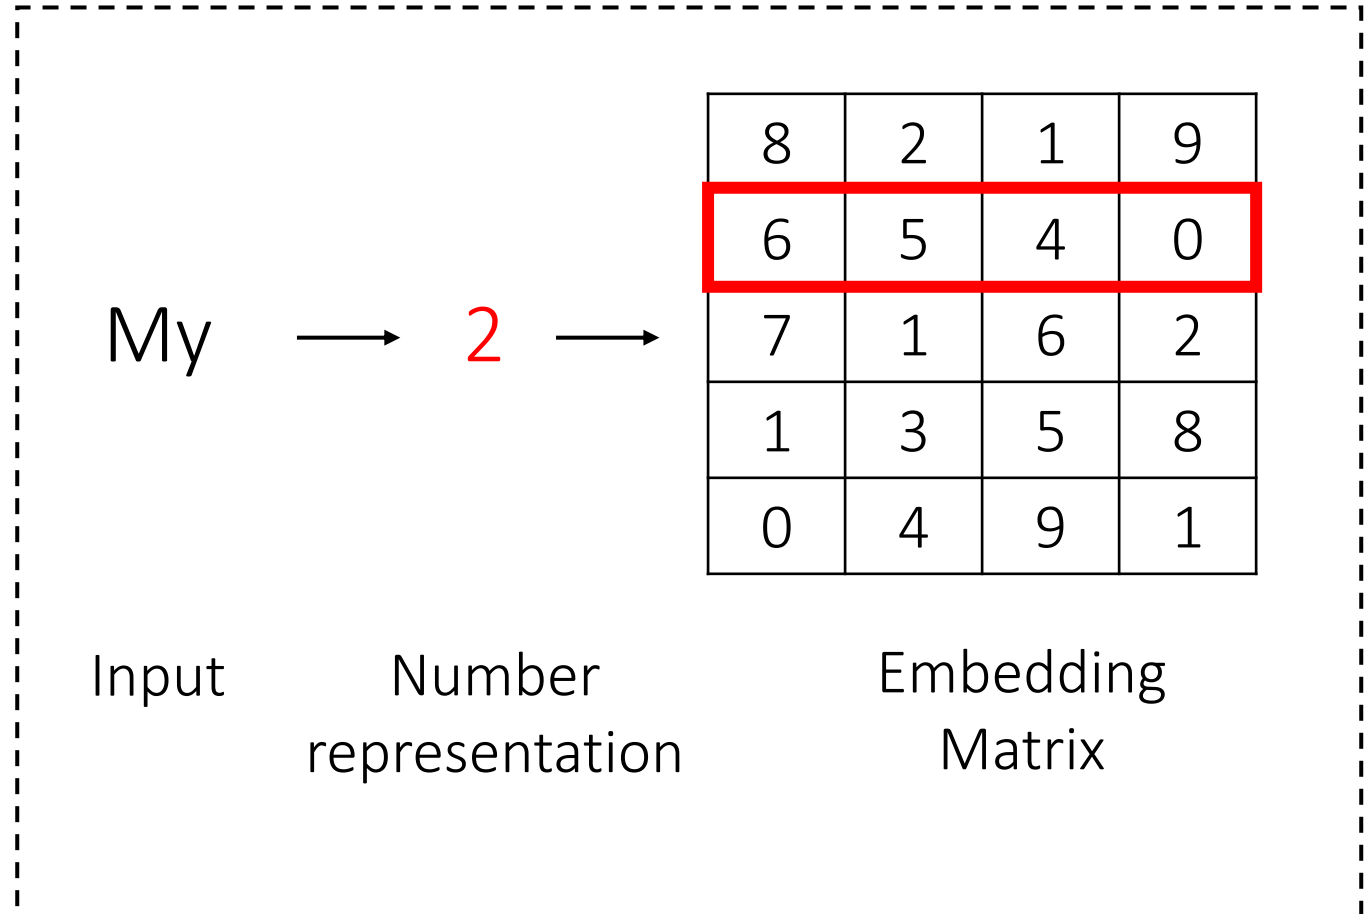



# Embedding



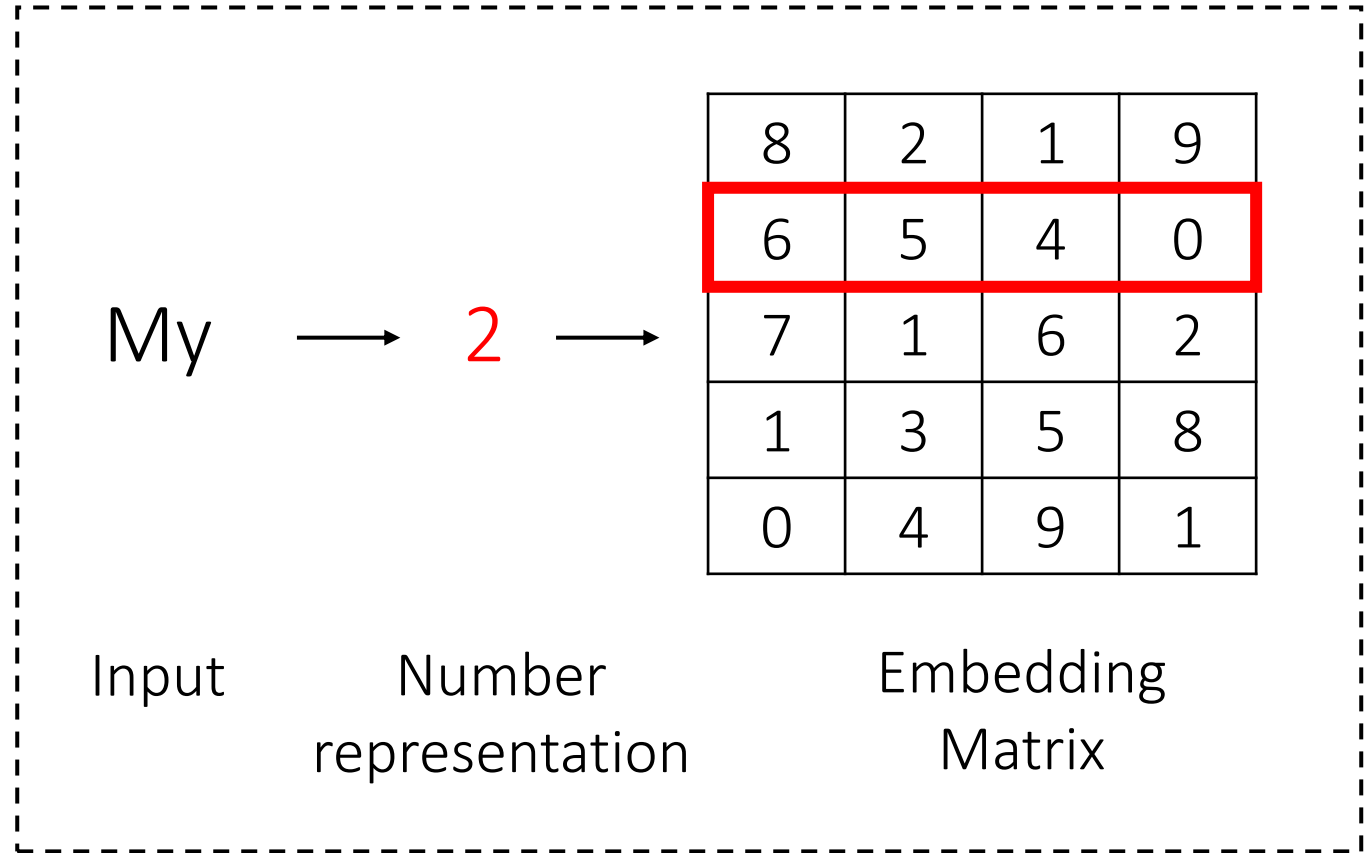



# Embedding





# Embedding



Embedding matrix is trainable



# Embedding

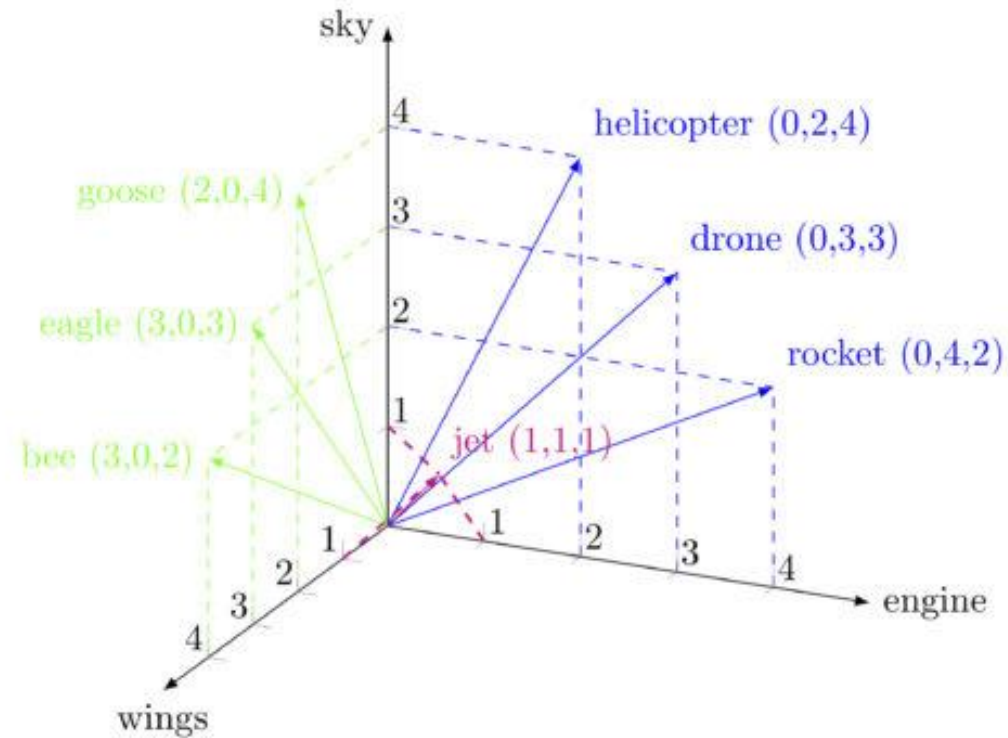


Image credit: hypotheses

Embedding matrix can encode **semantic representations** of inputs






# Backpropagation in RNNs

→ Forward  
← Backward

output

hidden

input





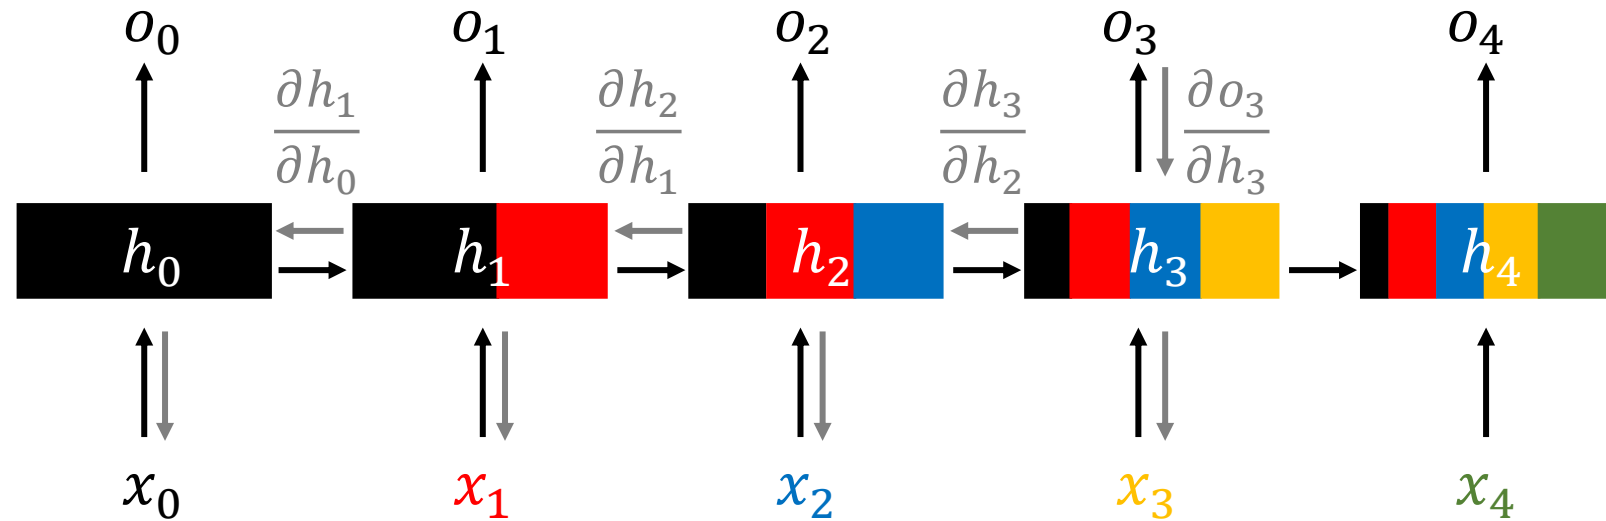
# Backpropagation in RNNs

→ Forward  
← Backward

output

hidden

input



Backpropagation is performed backward in time

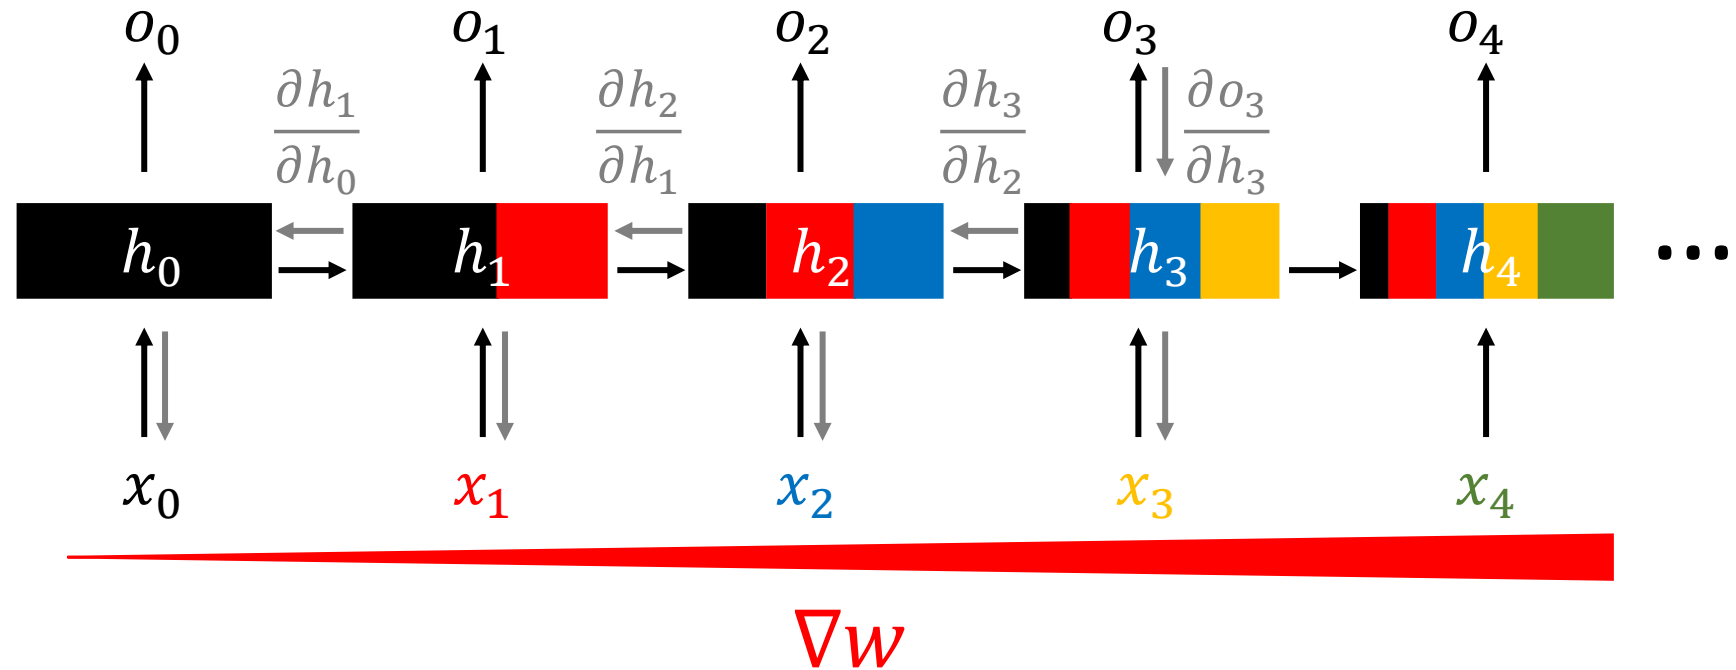
# Vanishing and Exploding Gradients

→ Forward  
← Backward

output

hidden

input






# Vanishing and Exploding Gradients

→ Forward  
← Backward

output

hidden

input



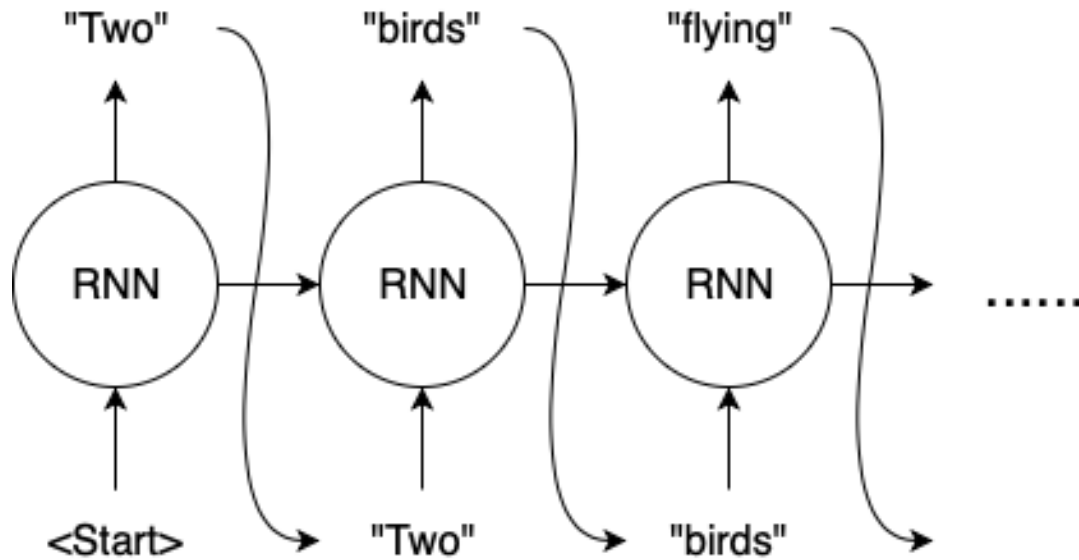
Longer input sequence →  
higher risk of Vanishing/Exploding Gradients!



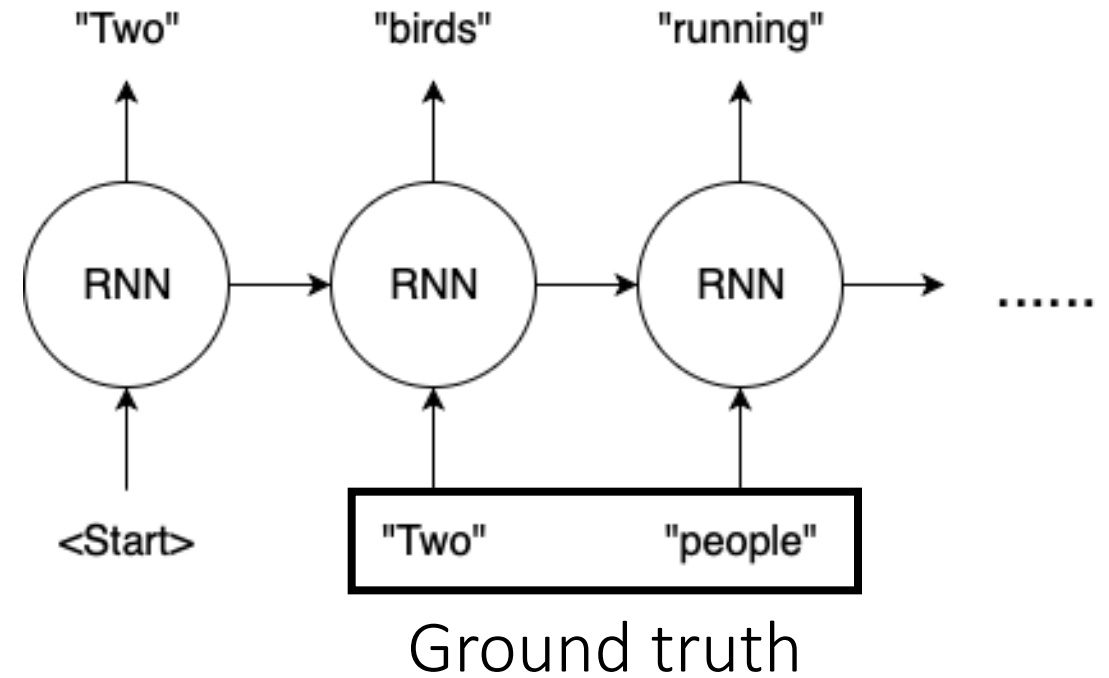
# Vanishing and Exploding Gradients

- Use gated RNN architecture e.g., LSTM, GRU (Next week)
- ReLU activation as nonlinearity
- Smaller number of sequence
- Smaller learning rate

# Training RNN with Teacher Forcing



Without Teacher Forcing



With Teacher Forcing



# RNN PROBLEM TYPES


RNN Configurations

RNN Extensions




# RNN Configurations


one to one




one to many




many to one



many to many




many to many





# One to One


one to one





# One to One

one to one




Identical to Feed Forward Network



# One to Many

one to many







# One to Many






Image captioning




Music generation



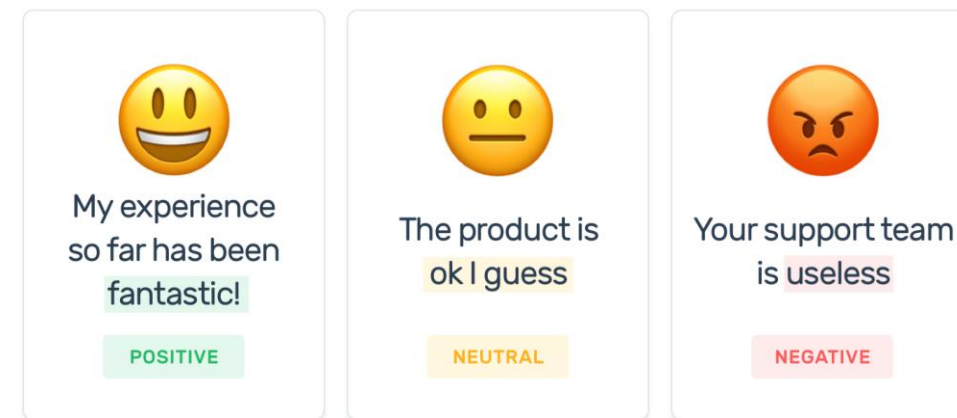
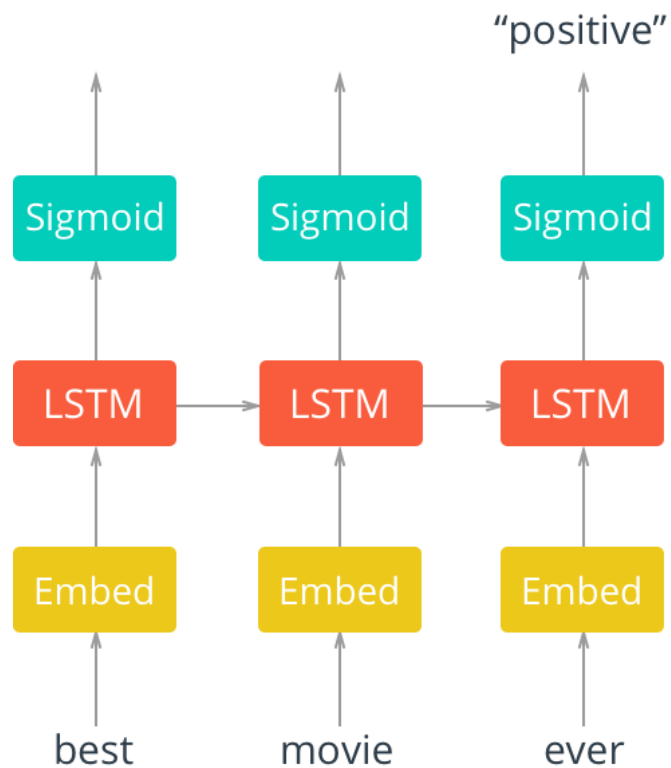



# Many to One





# Many to One




Sentiment Analysis




# Many to Many

many to many




many to many







# Many to Many

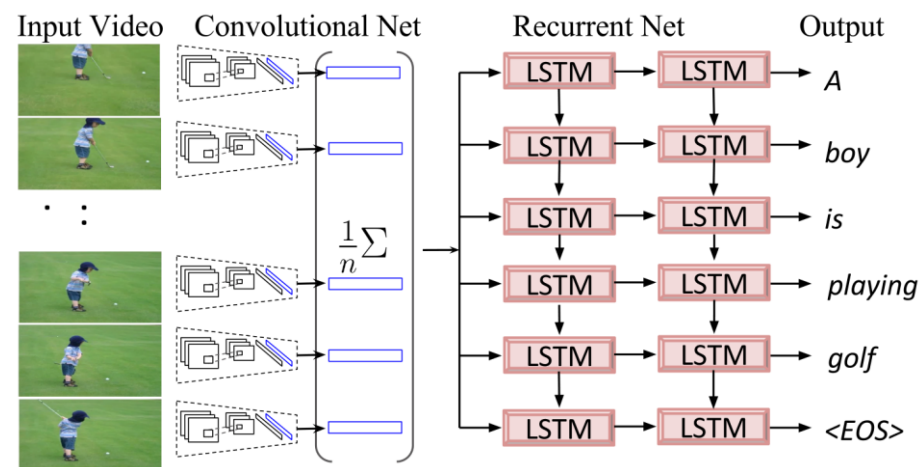
many to many



many to many




## Machine Translation




## Video Captioning




# RNN Extensions



Regular RNN




Deep RNN



Bi-directional RNN



# Deep RNN



Deep RNN

(+)


Can provide better performance  
Often used for complex problems

(-)

Potential for overfitting  
Longer training time



# Bi-directional RNNs



Bi-directional RNN

(+)

Higher performance in Natural Language Processing tasks

Suitable when both left and right contexts are used

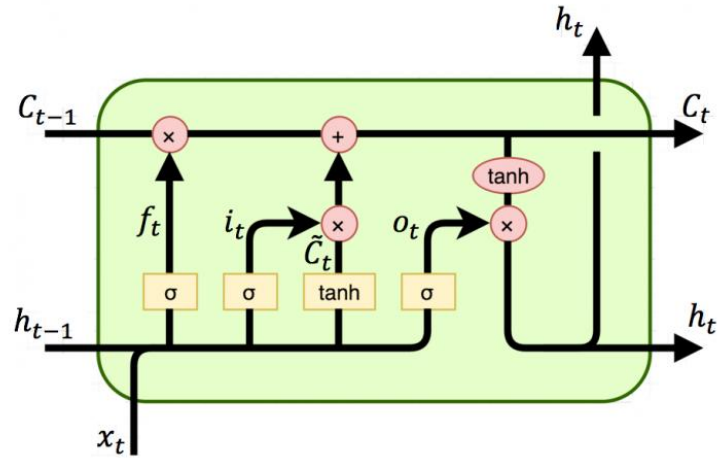
(-)

Harder to train than Uni-directional RNN  
Not suitable for real-time processing






# Next episode in EEP 596...



Long-short-term memory  
(LSTM)



Gated recurrent units  
(GRU)



# Next episode in EEP 596...

